

CprE 381: Computer Organization and Assembly Level Programming

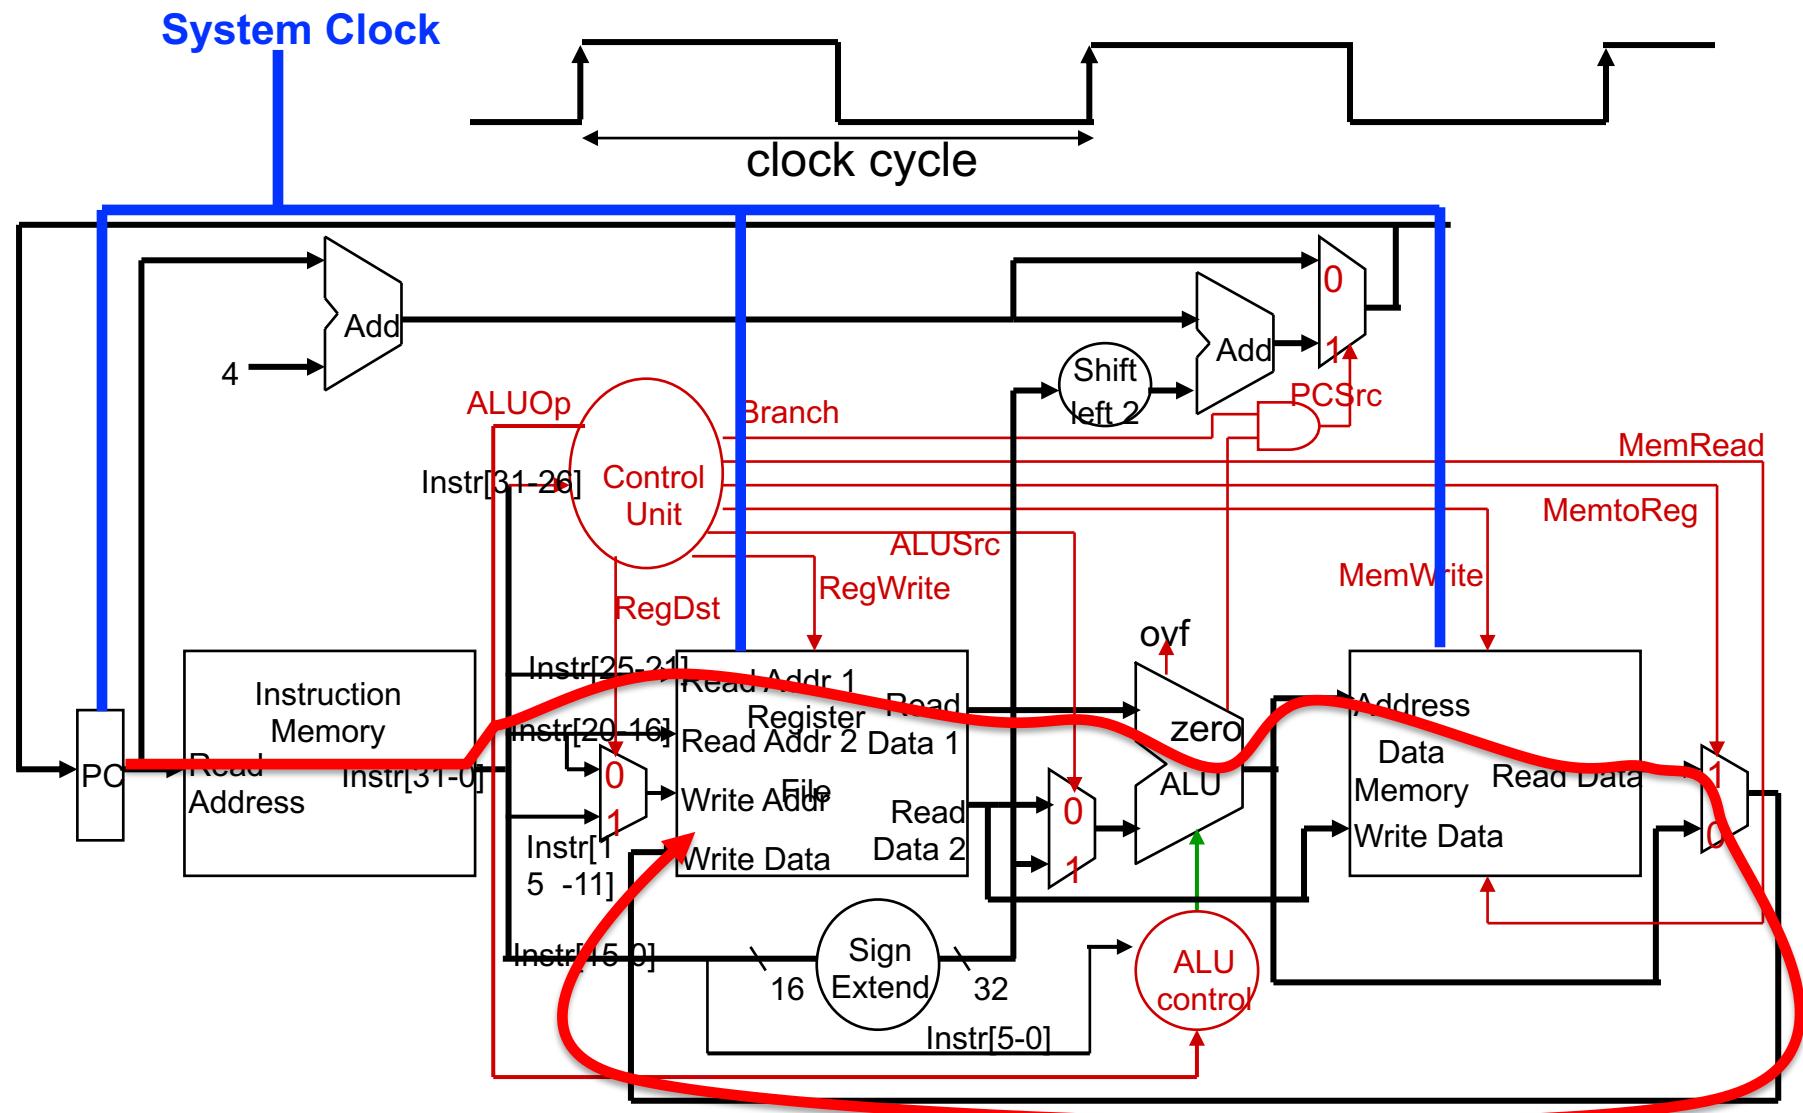
Pipelining

Henry Duwe
Electrical and Computer Engineering
Iowa State University

Administrative

- HW6 posted due Wed Mar 13
 - Make sure you have the new one
- Term Project
 - Part 2a due BEFORE Spring Break

Review: Worst Case Timing (Load Instruction)



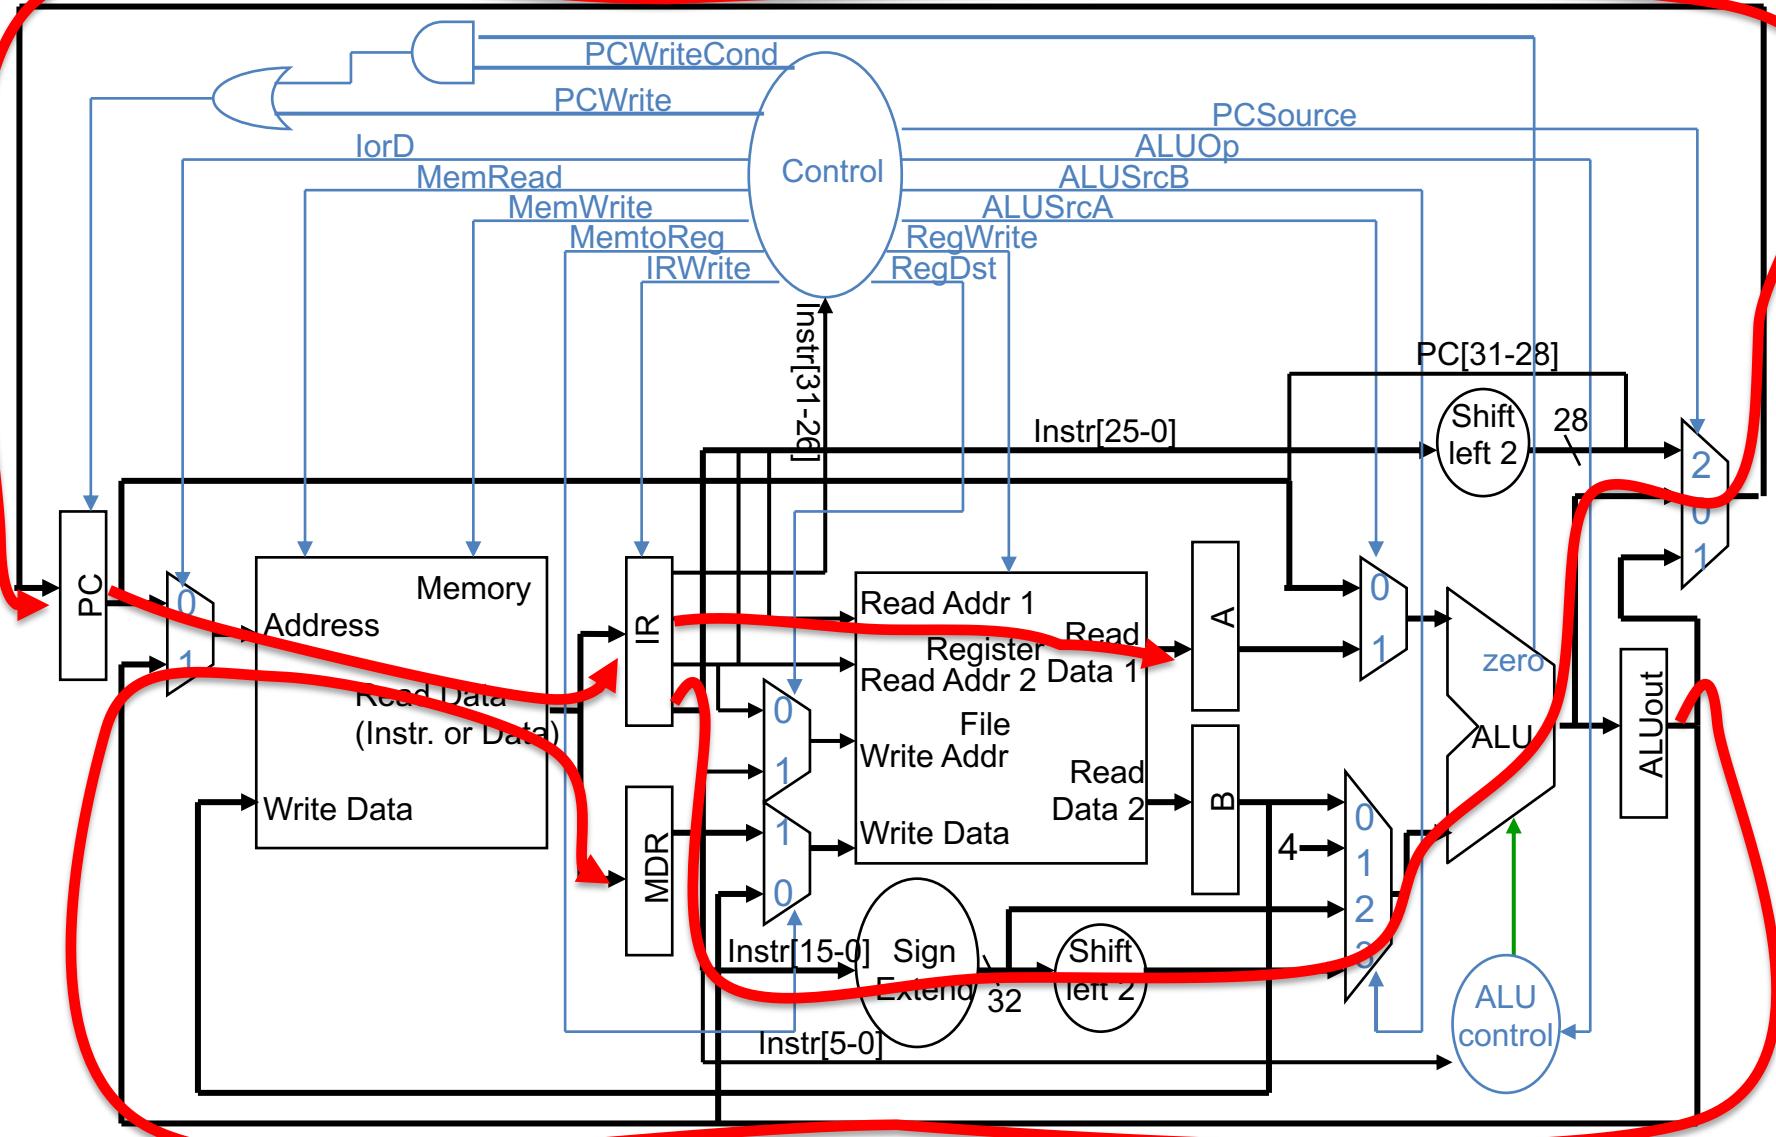
Review: Execution Time

- Drawing on the previous equation:

$$\text{Execution Time} = \# \text{ Instructions} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$$

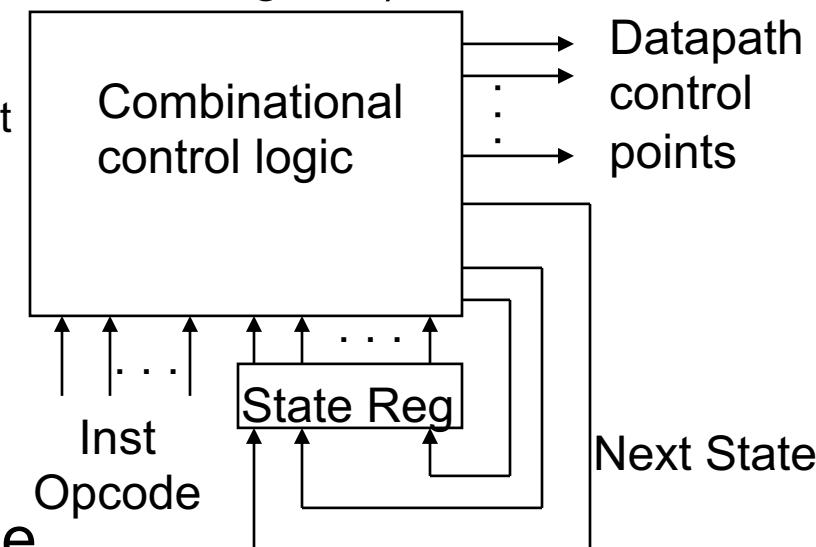
- To improve performance (i.e., reduce execution time)
 - Increase clock rate (decrease clock cycle time) OR
 - Decrease CPI OR
 - Reduce the number of instructions
- Designers balance cycle time against the number of cycles required
 - Improving one factor may make the other one worse...

Review: Multicycle Processor



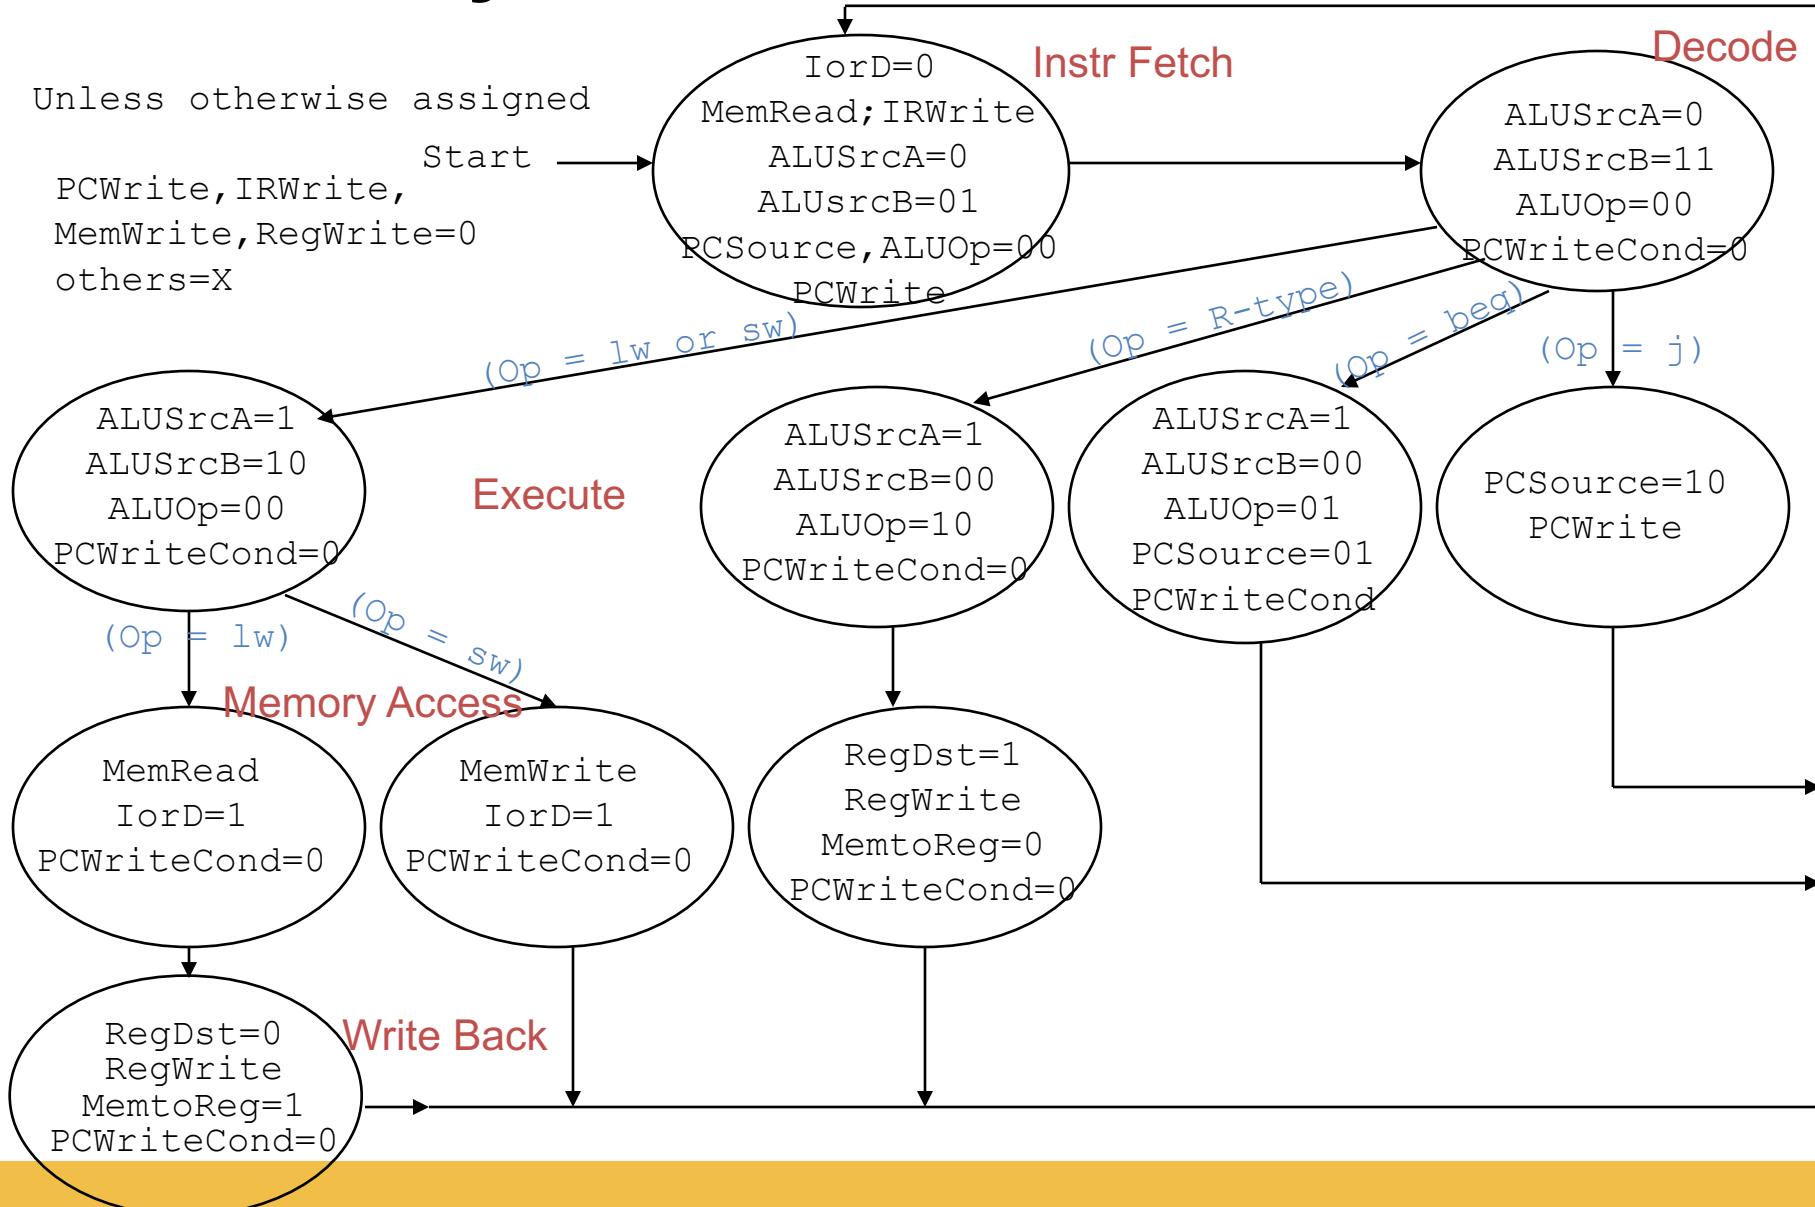
Review: Multicycle Control

- Multicycle datapath control signals are not determined solely by the bits in the instruction
 - e.g., op code bits tell what operation the ALU should be doing, but *not* what instruction cycle is to be done next
- We can use a finite state machine for control
 - A set of states (current state stored in State Register)
 - Next state function
 - determined by current state and the input
 - Output function
 - determined by current state)



- So we are using a **Moore** machine
(datapath control signals based only on current state)

Final Multicycle Control FSM



Simplifying the Control Unit Design

- For an implementation of the full MIPS ISA instructions can take from 3 clock cycles to 20+ clock cycles
 - Resulting in finite state machines with hundreds to thousands of states with even *more* arcs (state sequences)
 - Such state machine representations become impossibly complex
- Instead, can represent the set of control signals that are asserted during a state as a low-level control “instruction” to be executed by the datapath

Microinstructions

- “Executing” the microinstruction is equivalent to asserting the control signals specified by the microinstruction

Microprogramming (50k feet)

- A microinstruction has to specify
 - what control signals should be asserted
 - what microinstruction should be executed next
- Each microinstruction corresponds to one state in the FSM and is assigned a state number (or “address”)
 1. **Sequential** behavior – increment the state (address) of the current microinstruction to get to the state (address) of the next
 2. **Jump** to the microinstruction that begins execution of the next MIPS instruction (fetch state)
 3. **Branch** to a microinstruction based on control unit input using dispatch tables
 - need one for microinstructions following decode state
 - need another for microinstructions following memory address gen state
- The set of microinstructions that define a MIPS assembly language instruction (**macroinstruction**) is its **microroutine**

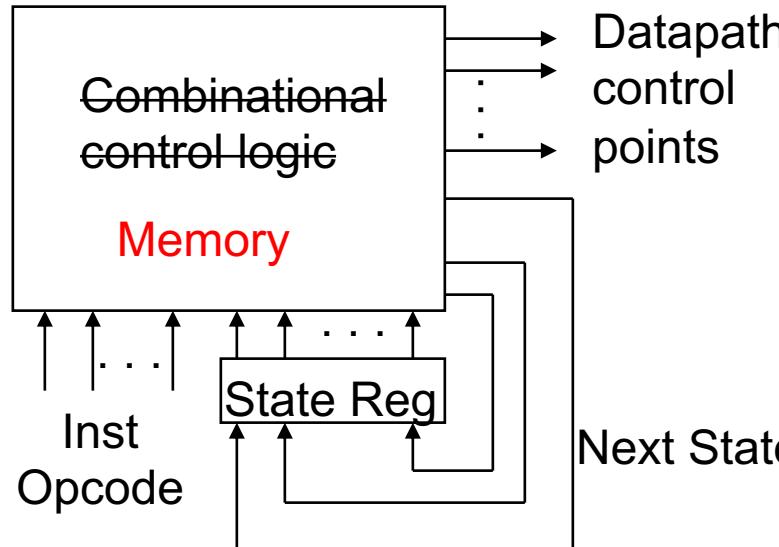
Microprogramming (50k feet)

- A microinstruction has to specify

- what
 - wh

- Each
FSM

1. Se cu
2. Ju M
3. Bi di
-



- The set of microinstructions that define a MIPS assembly language instruction (**macroinstruction**) is its **microroutine**

in the
ess")
of the
of the next
f the next
put using
address gen

Multicycle Advantages & Disadvantages

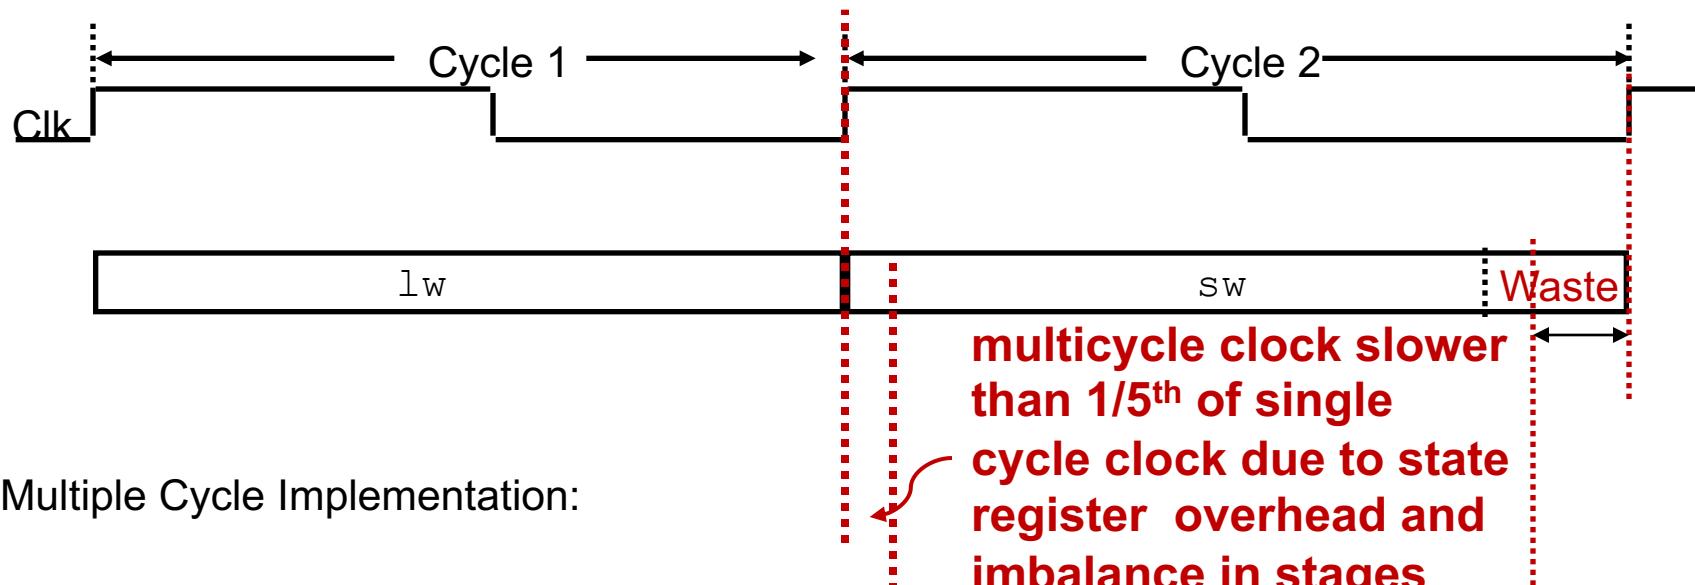
- Uses the clock cycle efficiently – the clock cycle is timed to accommodate the slowest instruction **step**
 - Balance the amount of work to be done in each step
 - Restrict each step to use only one major functional unit
- Multicycle implementations allow
 - Faster clock rates
 - Different instructions to take a different number of clock cycles
 - Functional units to be used more than once per instruction as long as they are used on different clock cycles

But

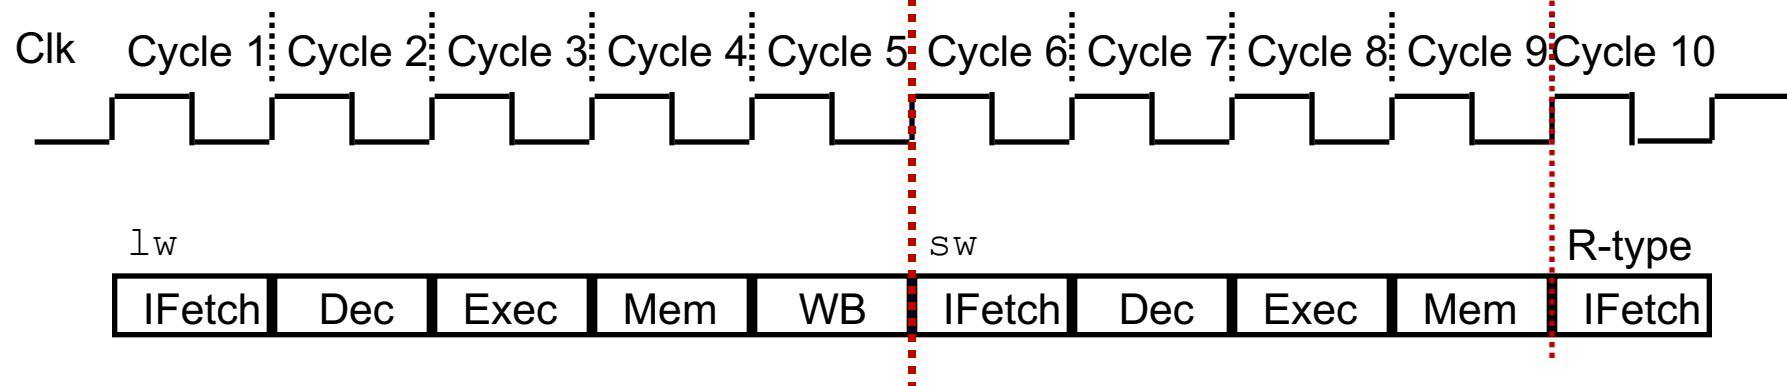
- Requires additional internal state registers, muxes, and more complicated (FSM) control
- **CPI > 1**

Single Cycle vs. Multiple Cycle Timing

Single Cycle Implementation:



Multiple Cycle Implementation:

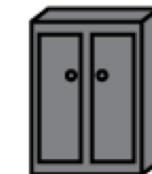


Increasing Parallelism

- Problem with the multi-cycle processor:
 - Each functional unit used once per instruction
 - Most of the time it is sitting waiting for its turn
- It's a circuit so it's calculating all the time, but it's waiting for valid data
 - Wasted time/energy
- Making instructions take more cycles can make machine faster!
 - Each instruction takes roughly the same time
 - While the CPI is much worse, the clock freq is much higher
 - **IDEA:** Overlap execution of multiple instructions at the same time
 - Different instructions will be active at the same time
 - This is called “Pipelining”
 - We will look at a 5 stage pipeline
 - Modern machines span the range from 3 (embedded micro-controllers) to 30+ (high performance processors) stages/instruction

Pipelining: You Do it All the Time!

- Ann, Brian, Cathy, Dave each have one load of clothes to wash, dry, fold, and put away
- Washer takes 30 minutes
- Dryer takes 30 minutes
- “Folder” takes 30 minutes
- “Stasher” takes 30 minutes to put clothes into drawers



Tech Update: Folder Exists!

 FoldiMate

[Home](#) [How it Works](#) [Speaking Technically](#) [FAQ](#) [Blog](#) [Check Status](#) [Want a FoldiMate?](#)

Speaking Technically...

 [RESTART](#)

 &  +  & 

AGE 6-XXL SIZE S-L [Learn More >](#)

 **Estimated Price: \$980**

We strive to make FoldiMate as affordable as we can. The estimate doesn't include tax and shipping, as regulations vary by region.

 **First Shipments:** Late 2019 [Why? >](#)

Starting in US then internationally based on demand.

 **Speed:** Under 4 minutes

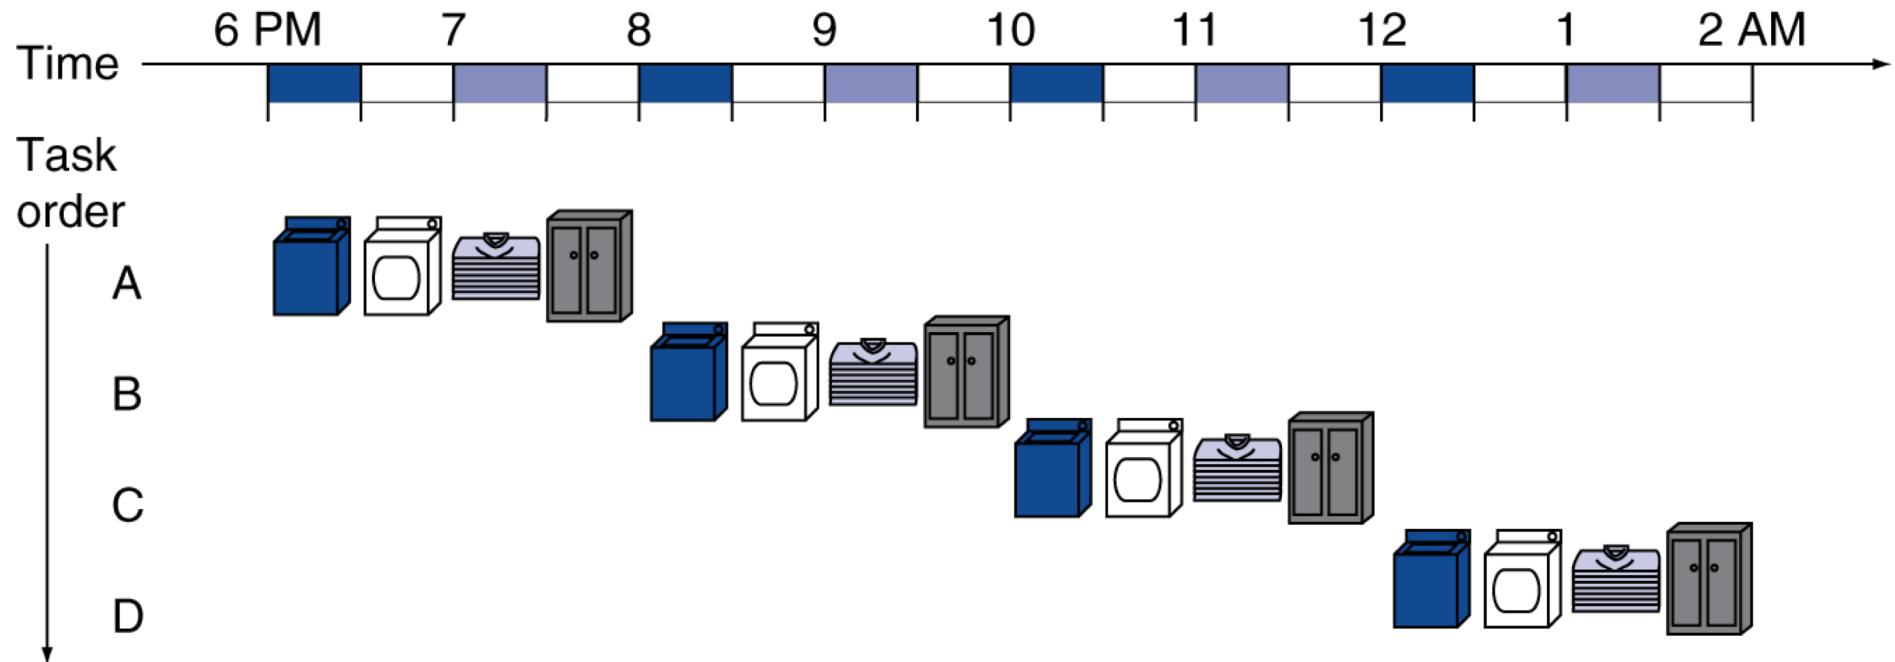
Imagine, getting your laundry folding done in less time than it takes to brew a cup of coffee!

 **Capacity:** Unlimited

FoldiMate folds as long as you continue clipping.

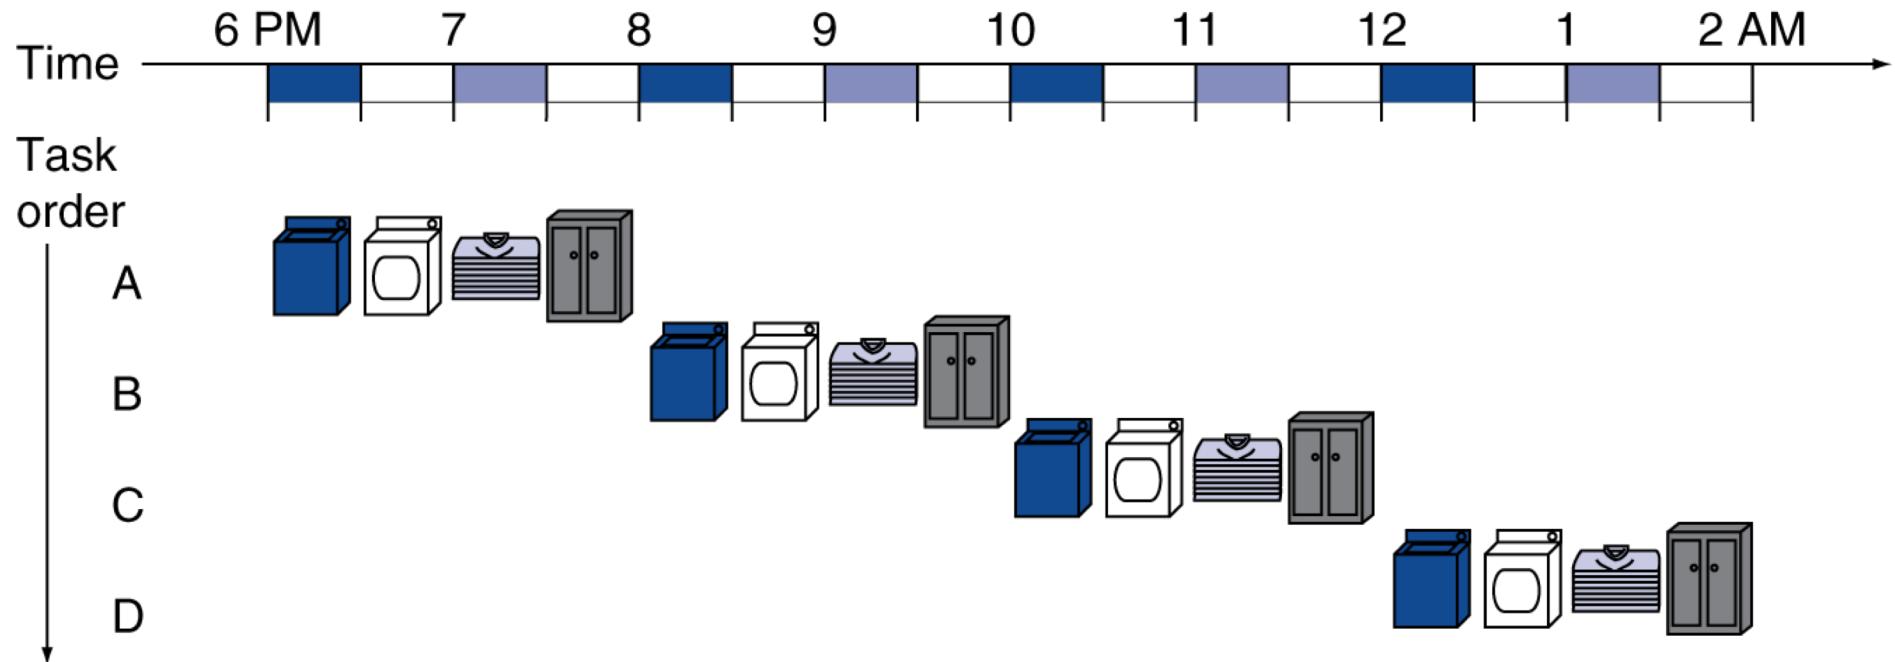


Sequential Laundry



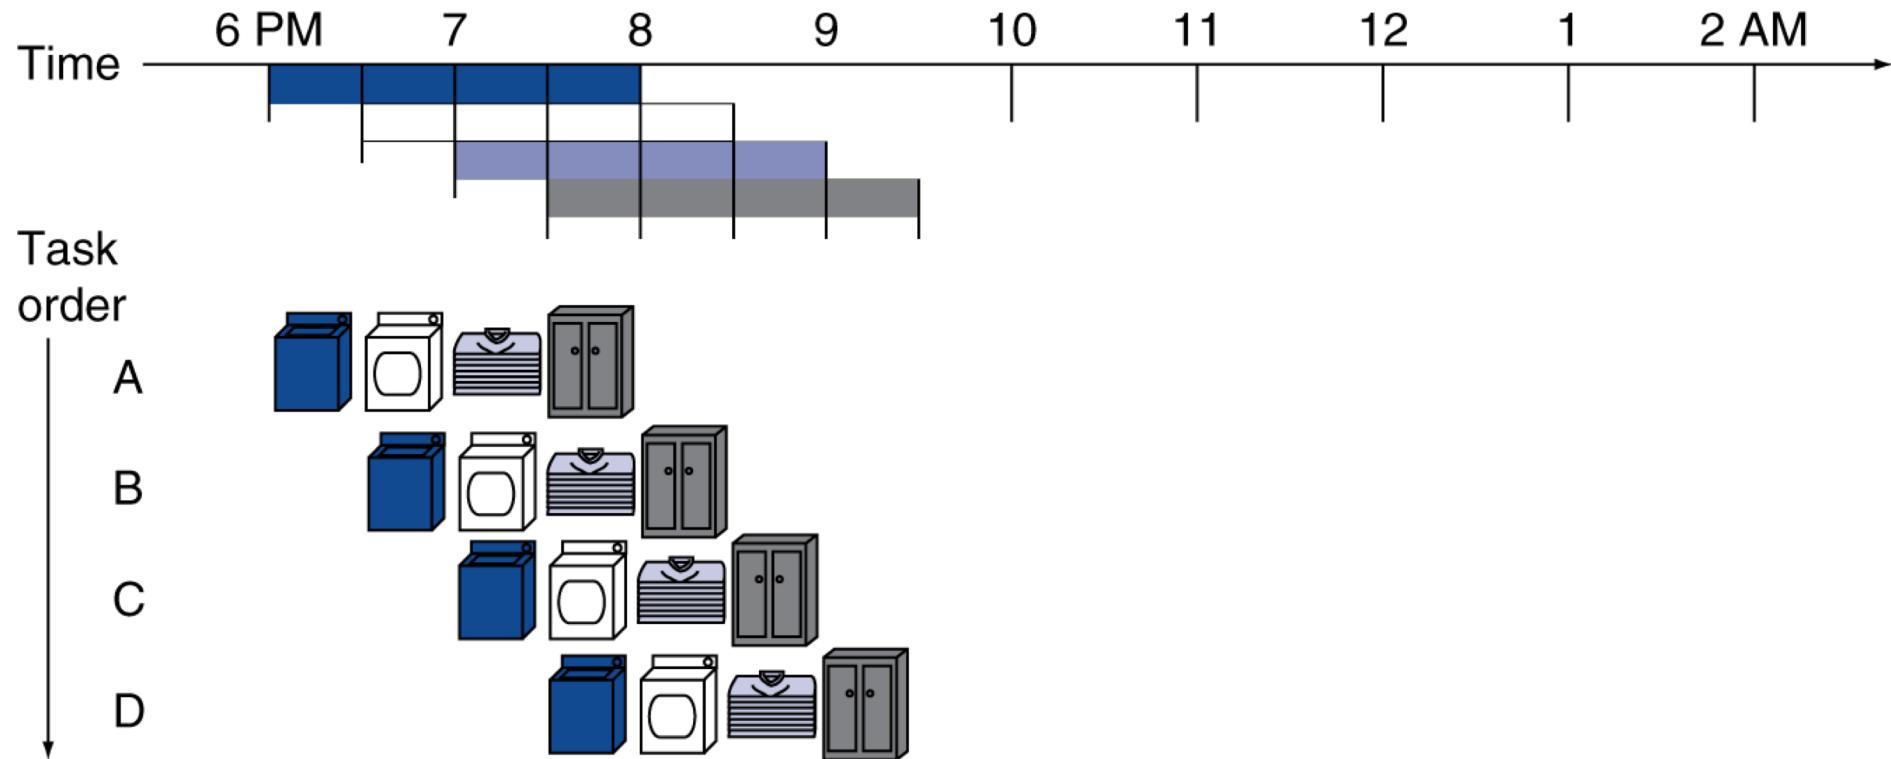
- Drawback: Sequential laundry takes 8 hours for 4 loads

Sequential Laundry



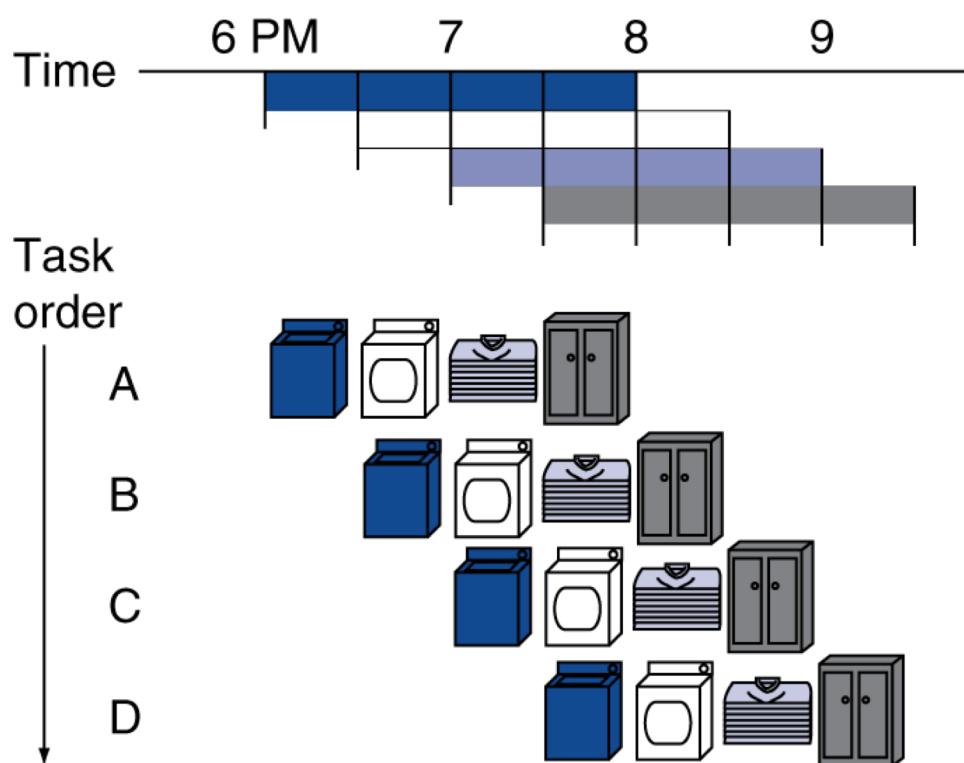
- Benefit: Allows Dave to work on 381 for 6 hours on Friday evening

Pipelined Laundry



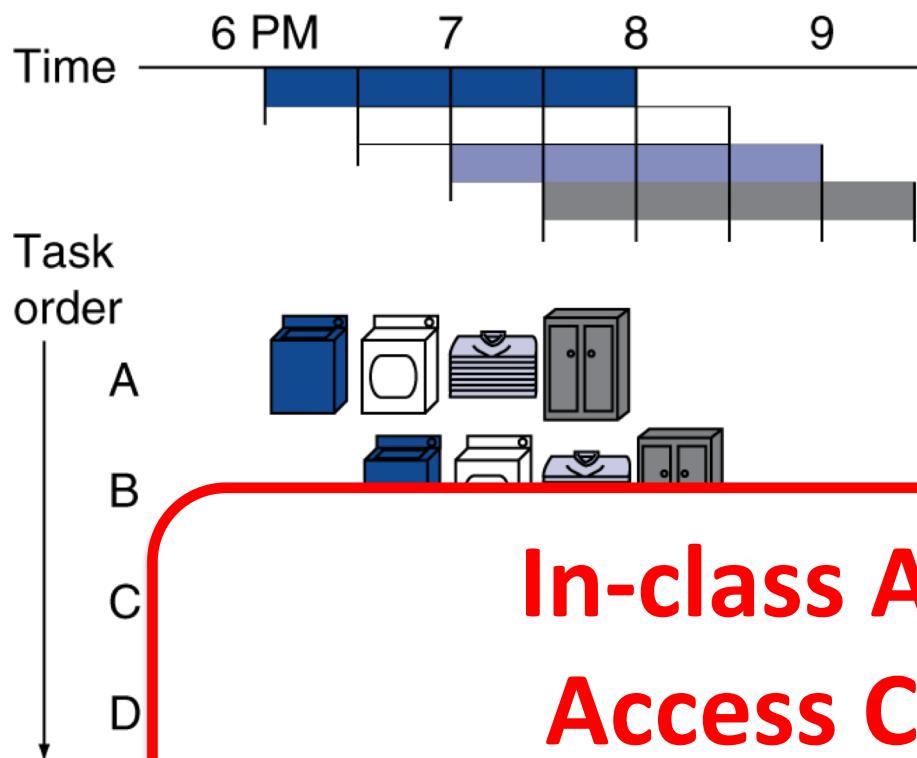
- Pipelined laundry takes 3.5 hours for 4 loads!
- Speedup = $8 / 3.5 = 2.3$
- General case: $0.5 * 4 * n / (0.5 * n + 1.5) \approx 4$ (# of stages)

Pipelining Lessons



- Pipelining doesn't help latency of single task, it helps throughput of entire workload
- Multiple tasks operating simultaneously using different resources
- Potential speedup = Number pipe stages
- Time to “fill” pipeline and time to “drain” it reduces speedup:
2.3X v. 4X in this example

Pipelining Lessons (cont.)



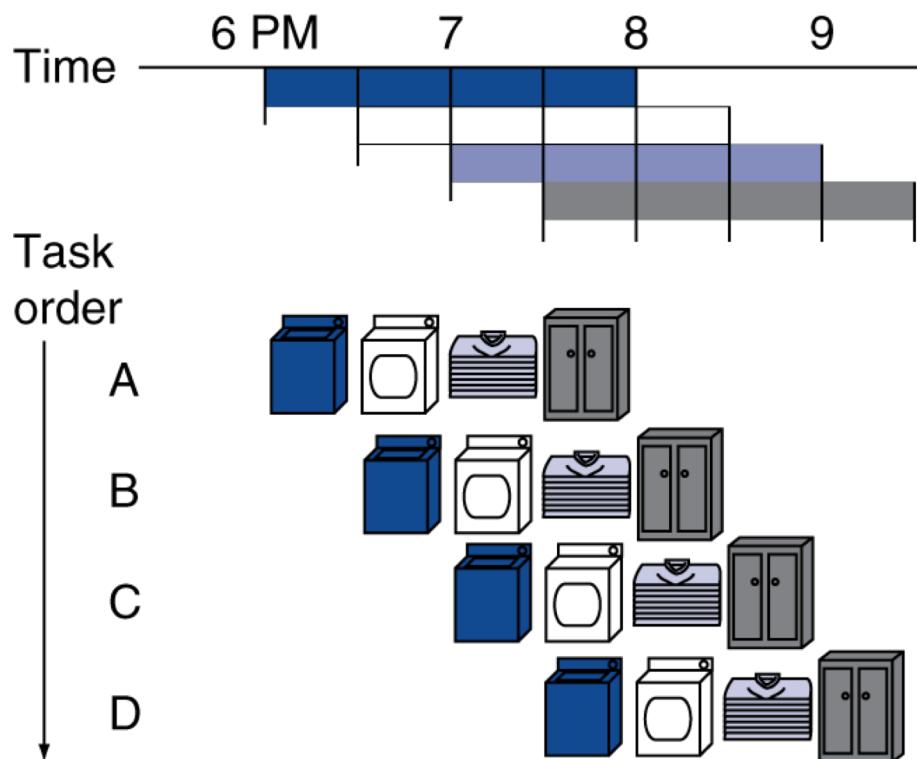
- Suppose new Washer takes 20 minutes, new Stasher takes 20 minutes. How much faster is pipeline?

In-class Assessment!

Access Code: iStack

Note: sharing access code to those outside of classroom or using access while outside of classroom is considered cheating

Pipelining Lessons (cont.)



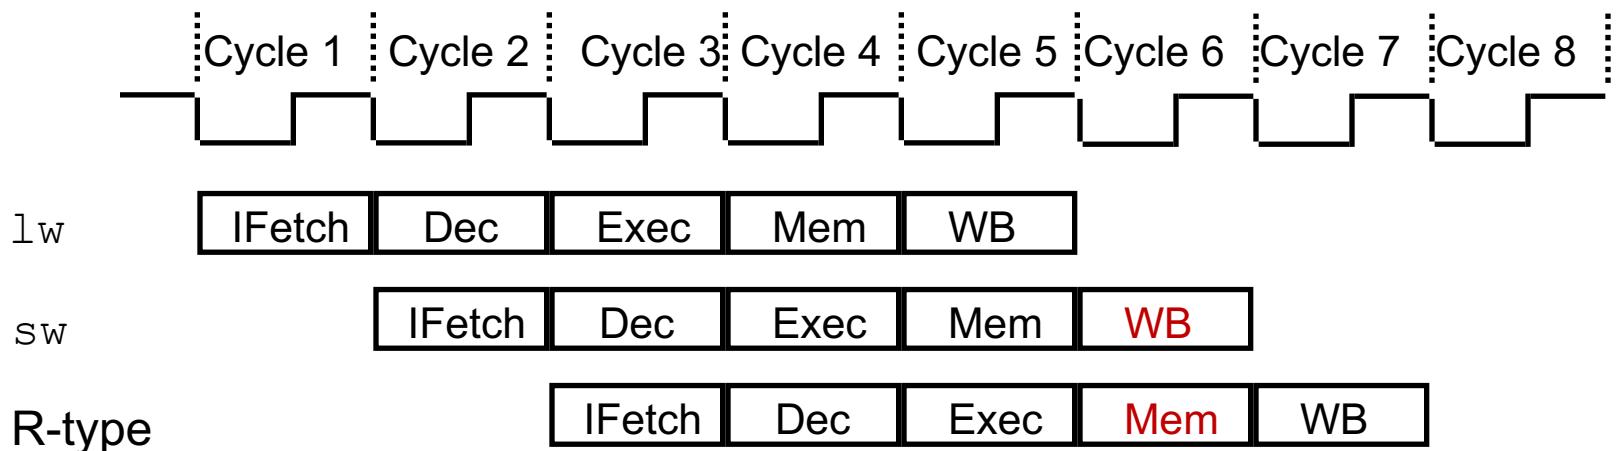
- Suppose new Washer takes 20 minutes, new Stasher takes 20 minutes. How much faster is pipeline?
- Pipeline rate limited by **slowest** pipeline stage
- Unbalanced lengths of pipe stages reduces speedup

Remember: 5 Stage MIPS Execution

- IF: Instruction Fetch
 - Fetch the instruction from memory
 - Increment the PC
- RF/ID: Register Fetch and Instruction Decode
 - Fetch base register
- EX: Execute
 - Calculate base + sign-extended offset
- MEM: Memory
 - Read the data from the data memory
- WB: Write back
 - Write the results back to the register file

A Pipelined MIPS Processor

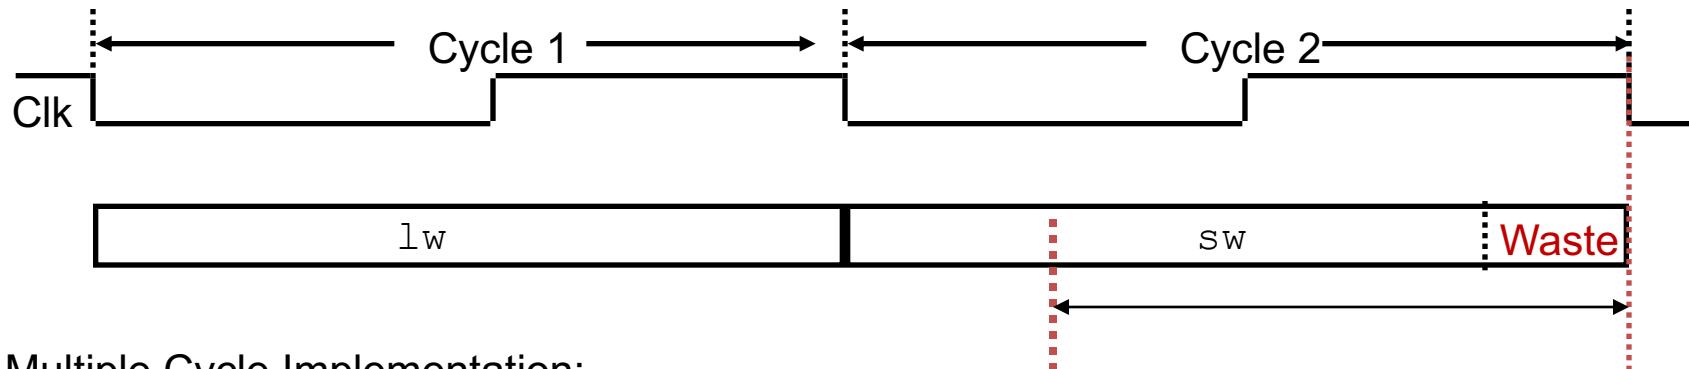
- Start the **next** instruction before the current one has completed
 - Improves **throughput** - total amount of work done in a given time
 - Instruction **latency** (execution time, delay time, response time - time from the start of an instruction to its completion) is *not* reduced



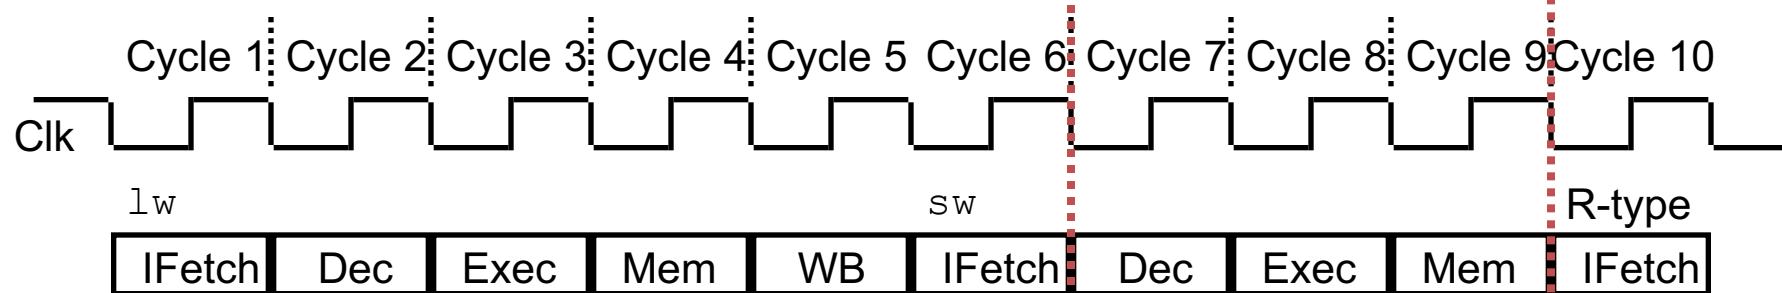
- Clock cycle (pipeline stage time) is limited by the slowest stage
- For some instructions, some stages are **wasted** cycles

Single Cycle, Multiple Cycle, vs. Pipeline

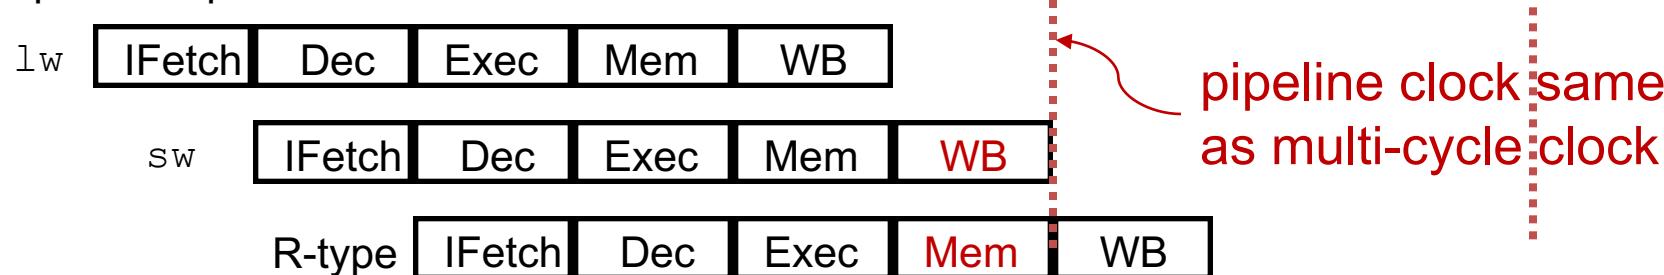
Single Cycle Implementation:



Multiple Cycle Implementation:

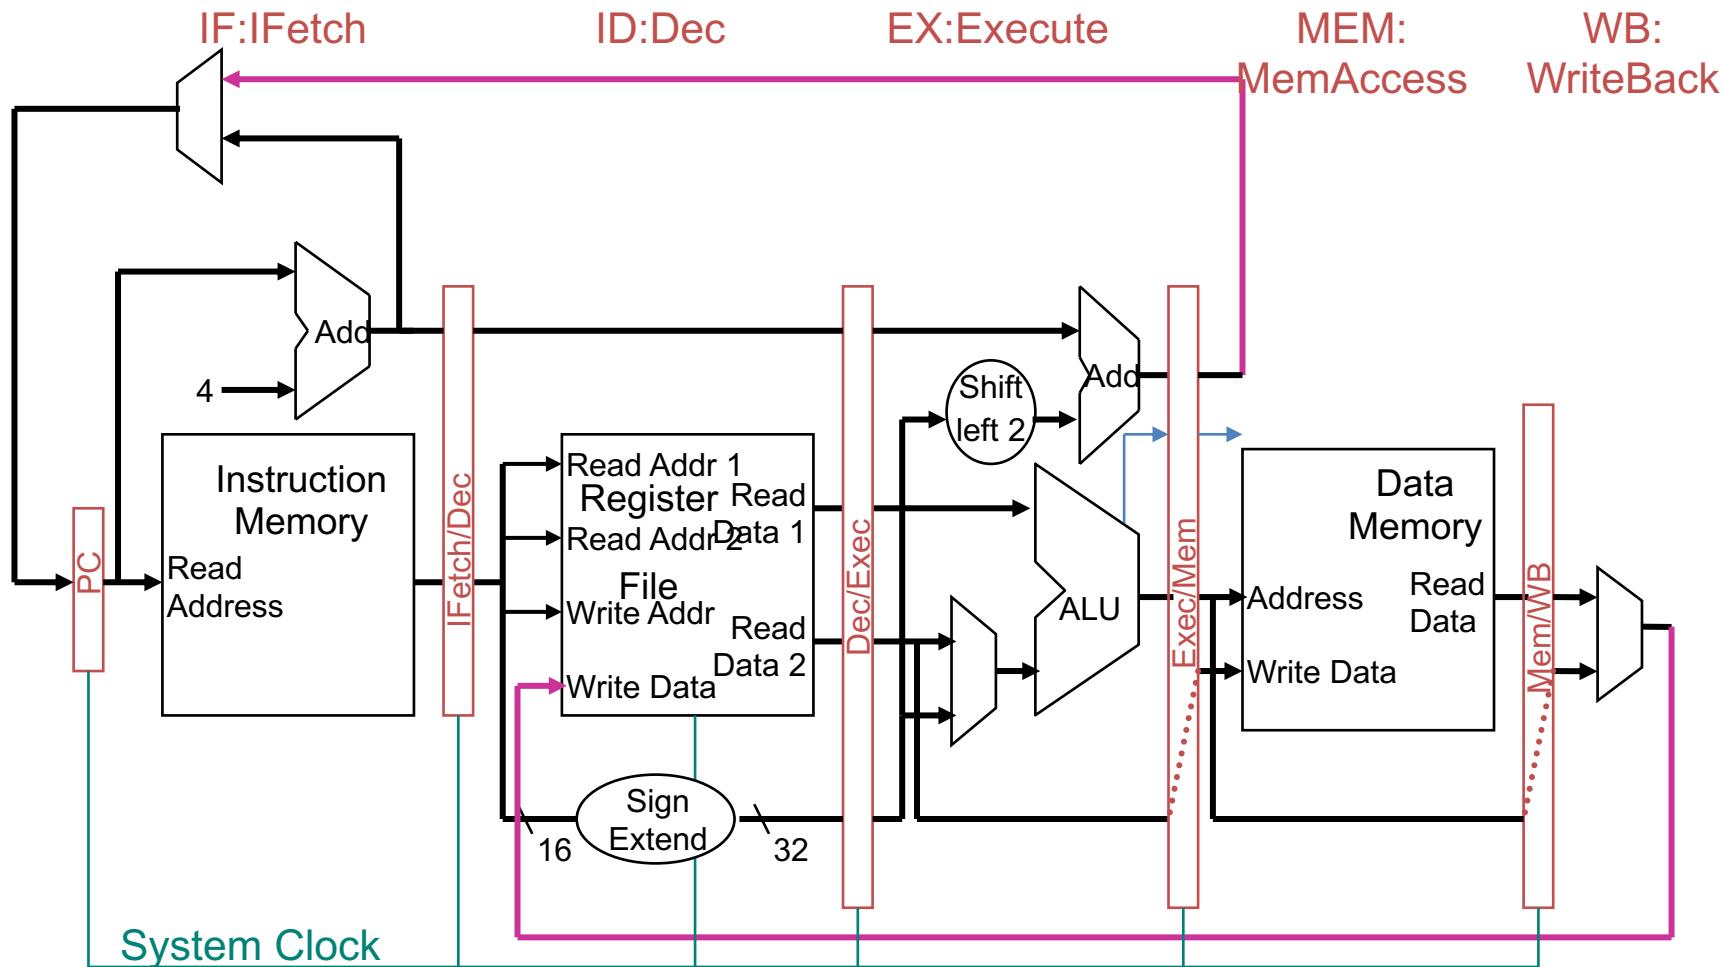


Pipeline Implementation:



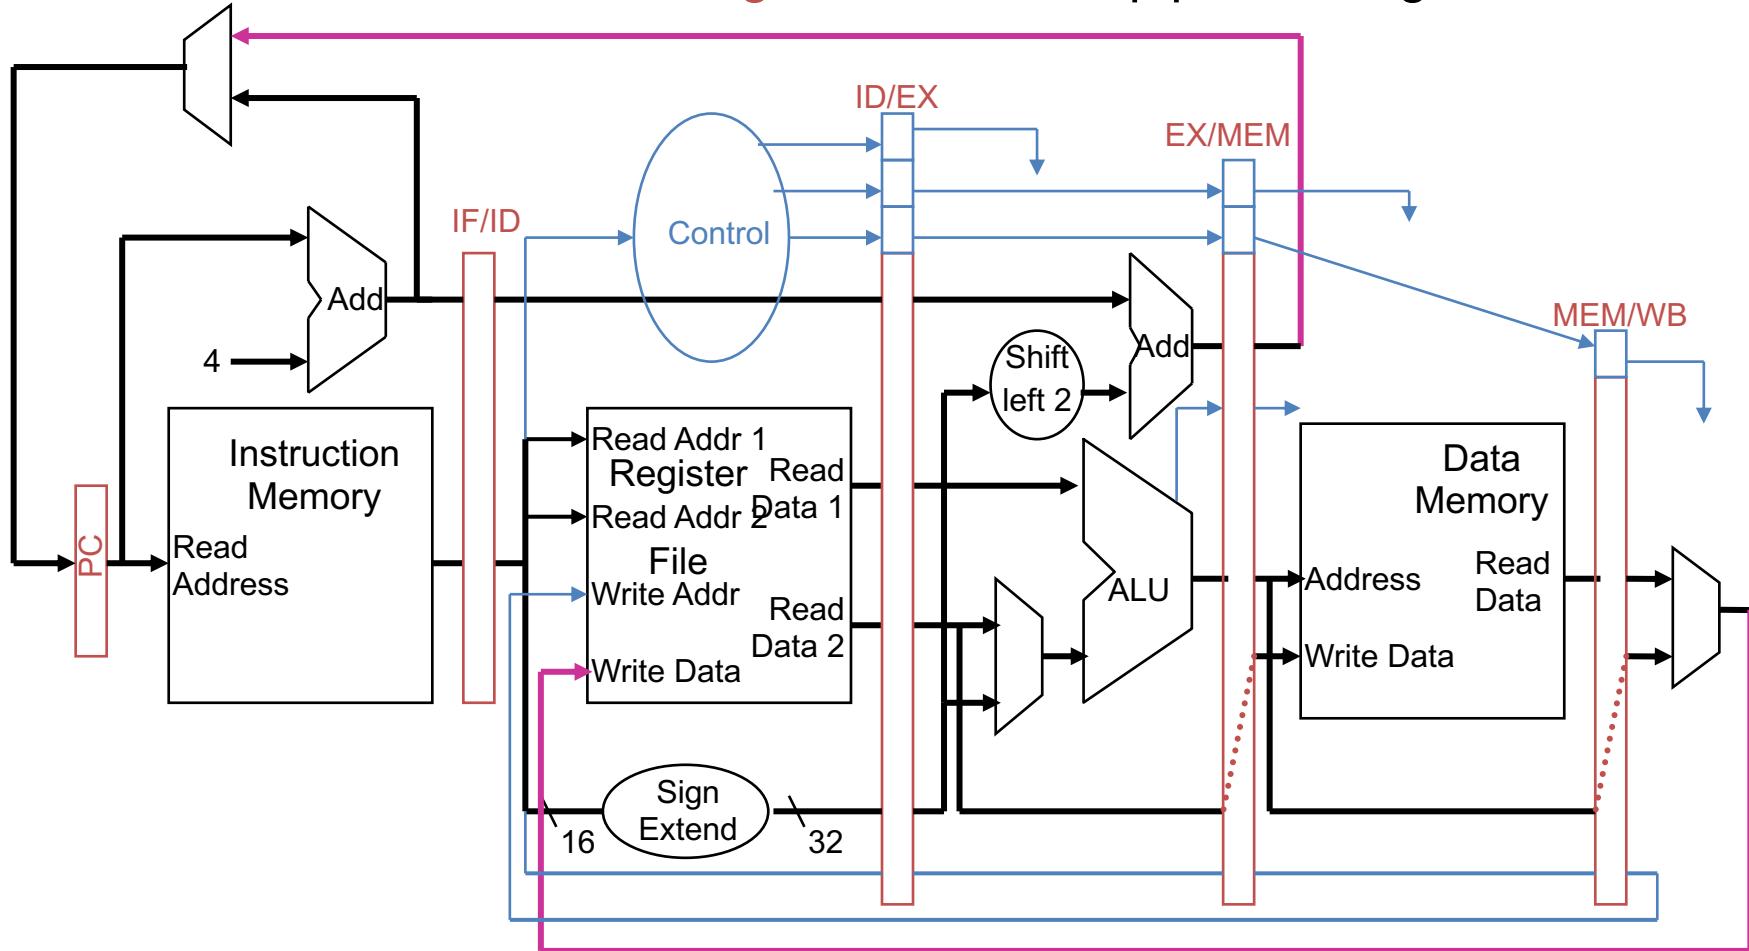
MIPS Pipeline Datapath Modifications

- What do we need to add/modify in our MIPS datapath?
 - State registers between each pipeline stage to **isolate** them



Preview: MIPS Pipeline Control Path Modifications

- All control signals can be determined during Decode
 - And held in the **state registers** between pipeline stages



Acknowledgments

- These slides contain material developed and copyright by:
 - Joe Zambreno (Iowa State)
 - David Patterson (UC Berkeley)
 - Mary Jane Irwin (Penn State)
 - Christos Kozyrakis (Stanford)
 - Onur Mutlu (Carnegie Mellon)
 - Krste Asanović (UC Berkeley)