

IOWA STATE UNIVERSITY

Department of Electrical and Computer Engineering

Lecture 04: OS Introduction III

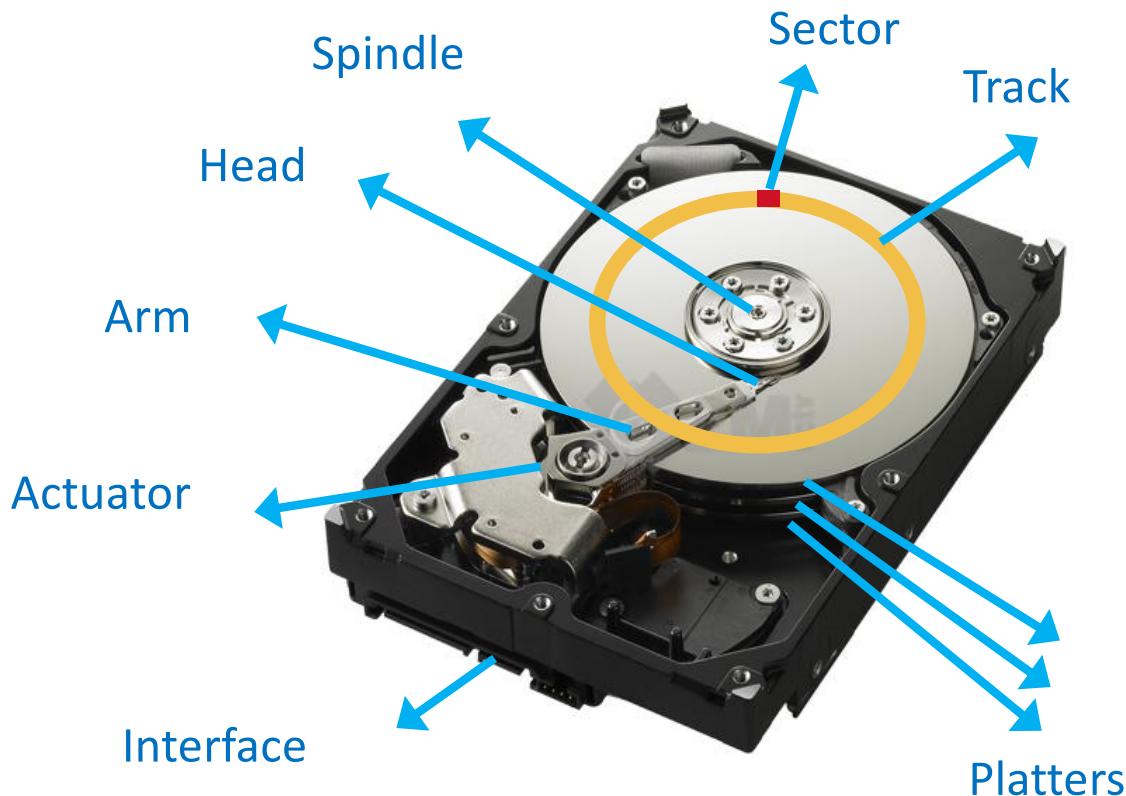


Agenda

- Recap
- OS Introduction III
 - OS Interface: System Calls
 - OS Internal Structures

Recap

- I/O device examples
 - hard disk drives (HDDs) & solid-state drives (SSDs)

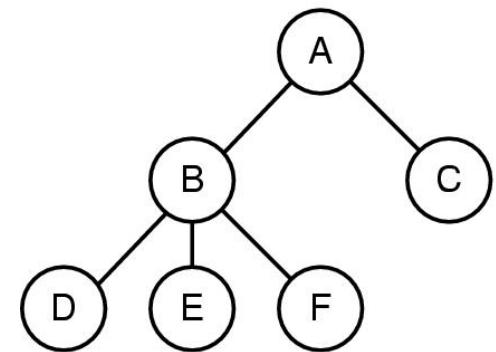


Recap

- Common OS Abstractions for HW
 - CPU
 - process and/or thread
 - Memory
 - address space
 - Disks
 - files

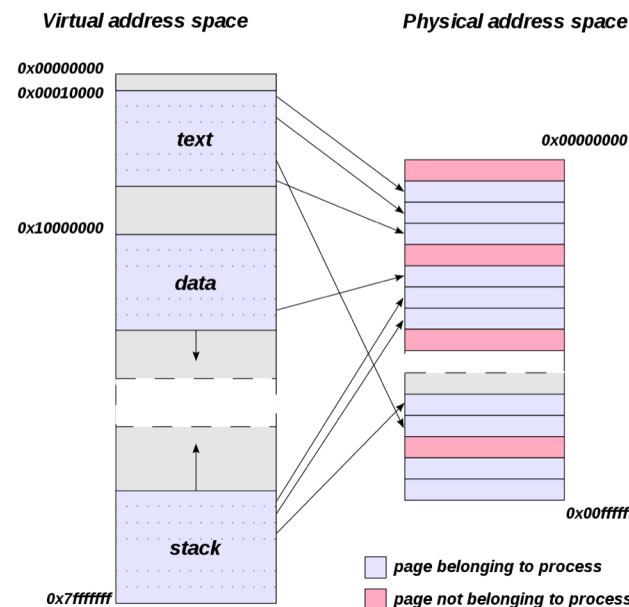
Recap

- Process
 - A program *in execution*
 - holds all the info needed to run a program
 - Address space
 - Program (text), data, stack
 - Register values
 - Program counter, stack pointer, etc
- Process Tree
 - a hierarchy of related processes
- Process management
 - create, schedule, suspend, resume, kill, communication, synchronization, ...



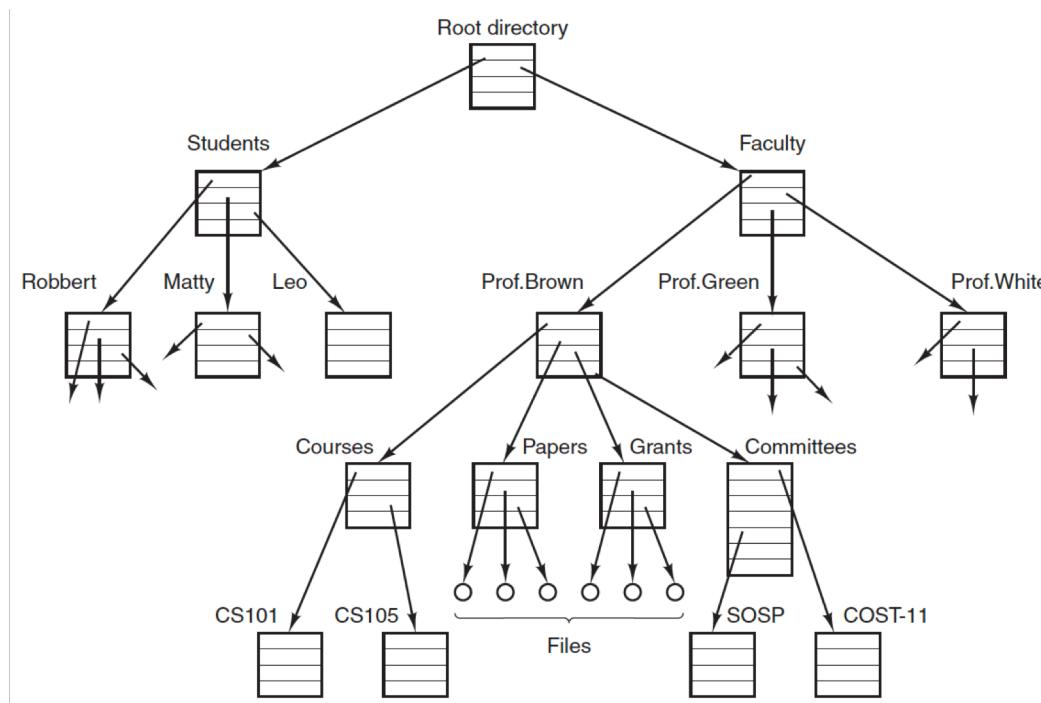
Recap

- Address space
 - A range of memory addresses accessed by the process
 - each process has its own **virtual address space**.
 - OS maps virtual address space onto the physical memory



Recap

- Files
 - “Everything is a file” on Linux
 - resources are represented as files in a file system tree
 - retrievable via a path from the root

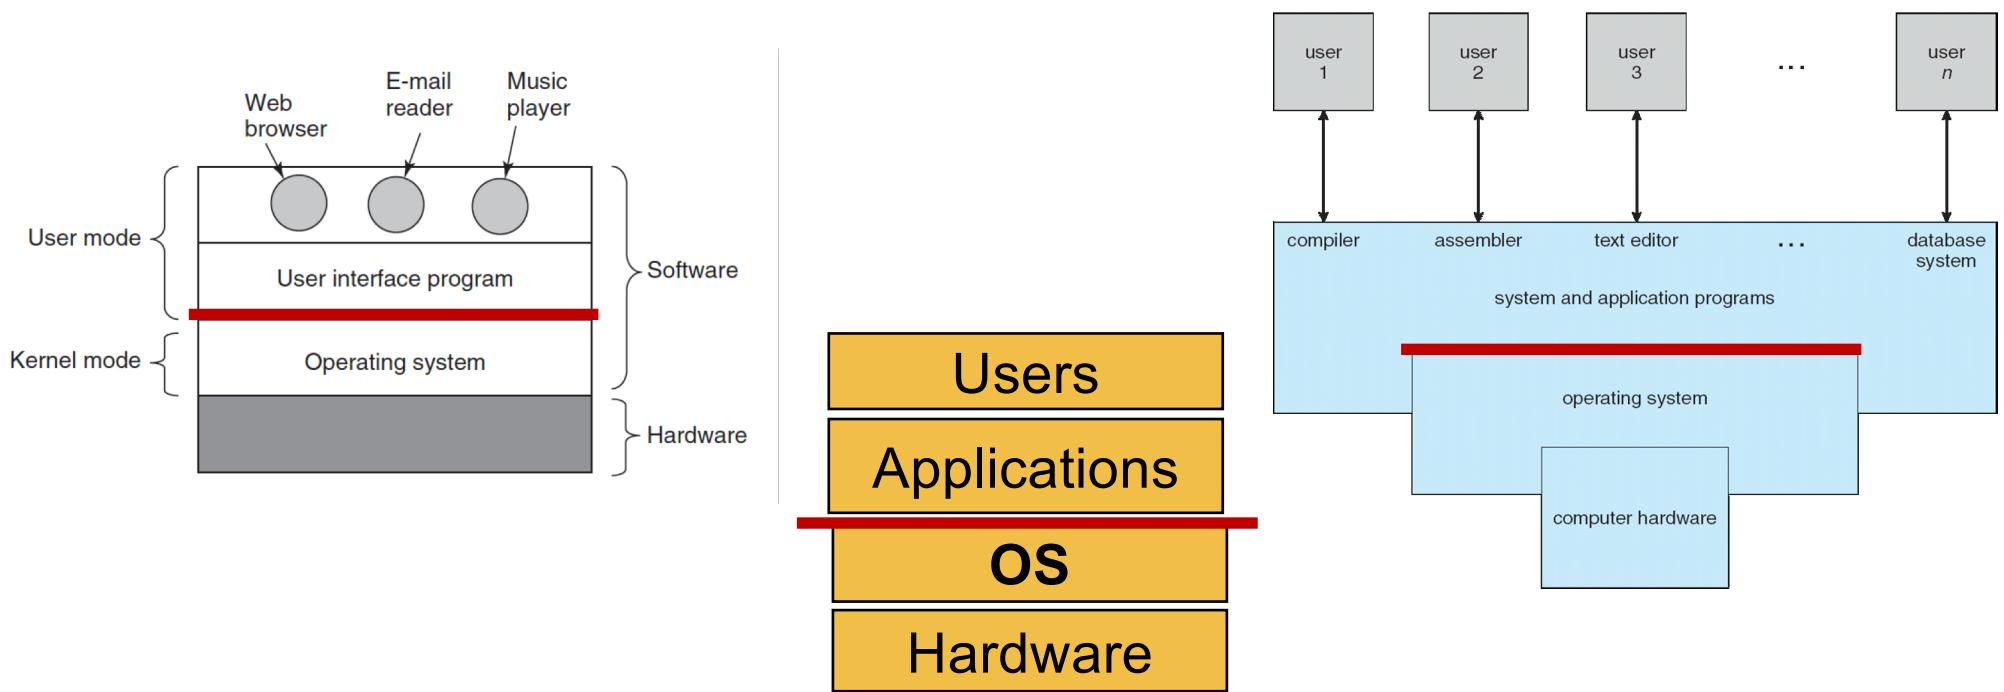


Agenda

- Recap
- OS Introduction III
 - OS Interface: System Calls
 - OS Internal Structures

System Calls

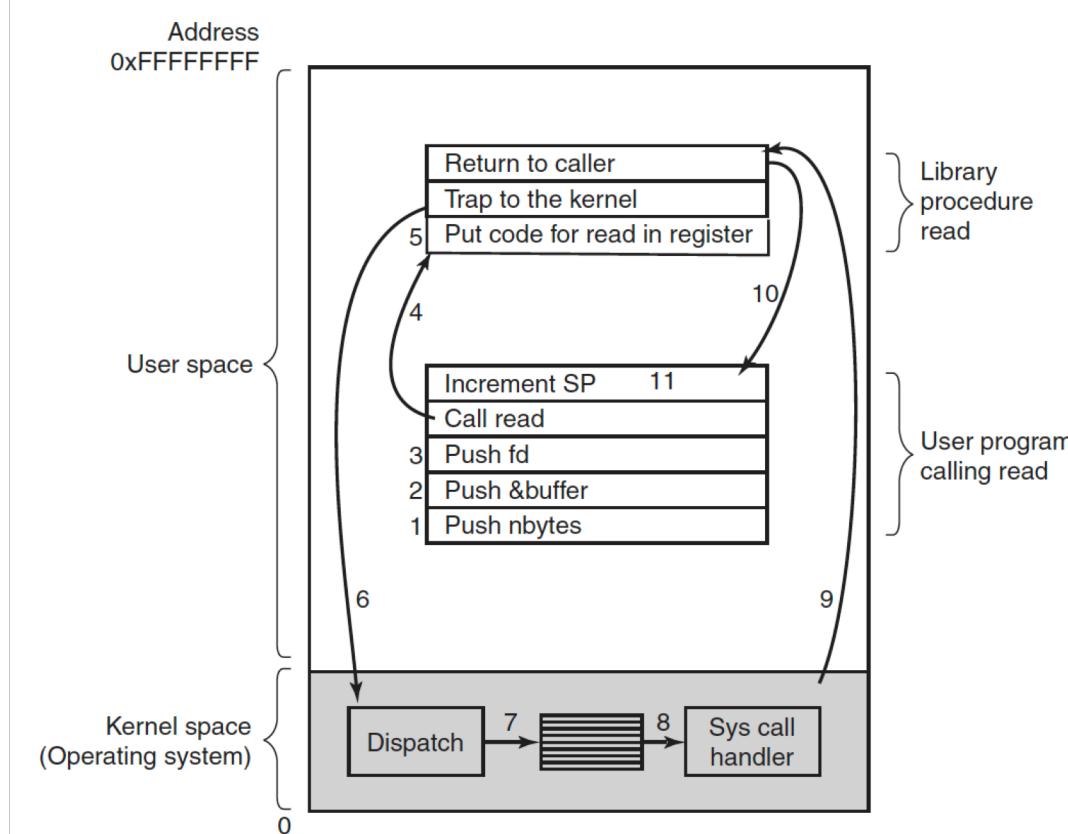
- Function calls implemented by OS
 - interface to apps/users
 - may be wrapped in library calls (e.g., glibc)
 - many follow the POSIX standard



System Calls

- Examples
 - a system call directly invoked by a program

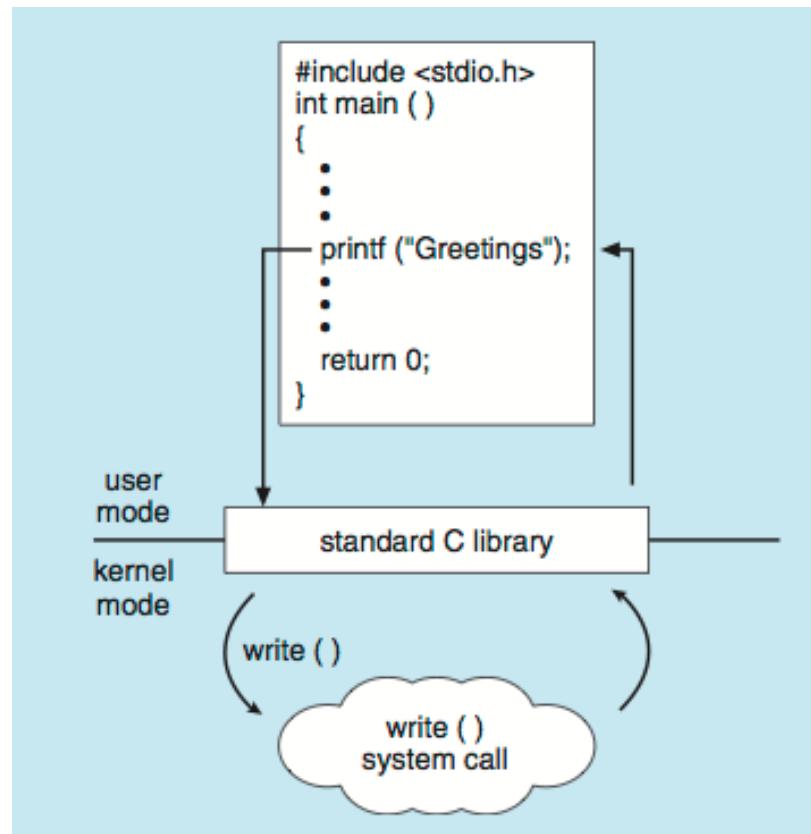
count = read(fd, buffer, nbytes);



System Calls

- Examples
 - a system call wrapped in a library function

```
ssize_t write(int fd, const void *buf, size_t count); /*Linux*/
```



System Calls

- Examples
 - process management

Call	Description
pid = fork()	Create a child process identical to the parent
pid = waitpid(pid, &statloc, options)	Wait for a child to terminate
s = execve(name, argv, environp)	Replace a process' core image
exit(status)	Terminate process execution and return status

System Calls

- Examples
 - process management
 - a simplified shell

```
#define TRUE 1

while (TRUE) {
    type_prompt();
    read_command(command, parameters);

    if (fork() != 0) {
        /* Parent code. */
        waitpid(-1, &status, 0);
    } else {
        /* Child code. */
        execve(command, parameters, 0);
    }
}

/* repeat forever */
/* display prompt on the screen */
/* read input from terminal */

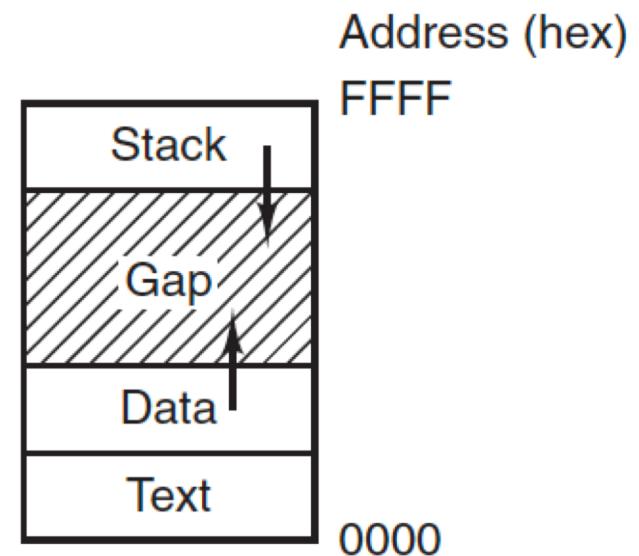
/* fork off child process */

/* wait for child to exit */

/* execute command */
```

System Calls

- Examples
 - UNIX processes have their memory divided into 3 segments
 - text: the program code
 - data: variables
 - extend explicitly via brk syscall
 - stack: for function calls



System Calls

- Examples
 - file/directory/file-system management

Call	Description
<code>fd = open(file, how, ...)</code>	Open a file for reading, writing, or both
<code>s = close(fd)</code>	Close an open file
<code>n = read(fd, buffer, nbytes)</code>	Read data from a file into a buffer
<code>n = write(fd, buffer, nbytes)</code>	Write data from a buffer into a file
<code>position = lseek(fd, offset, whence)</code>	Move the file pointer
<code>s = stat(name, &buf)</code>	Get a file's status information

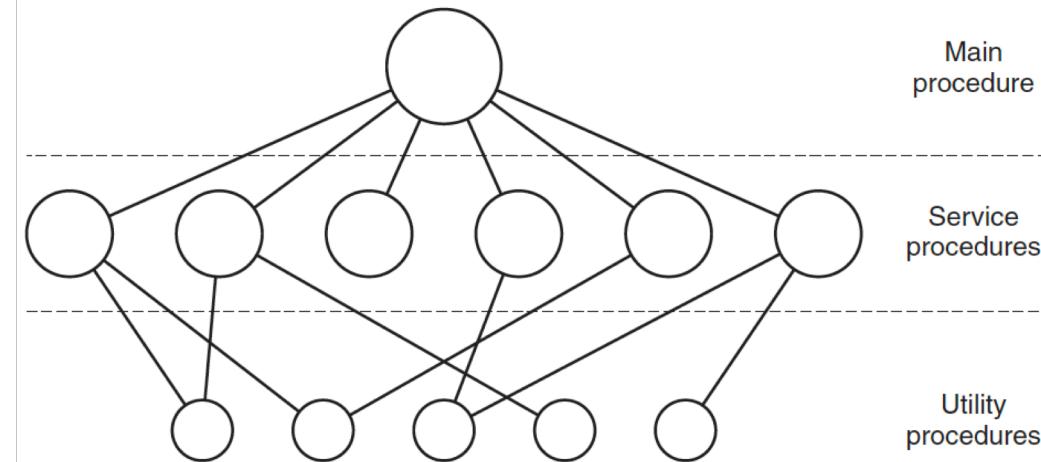
Call	Description
<code>s = mkdir(name, mode)</code>	Create a new directory
<code>s = rmdir(name)</code>	Remove an empty directory
<code>s = link(name1, name2)</code>	Create a new entry, name2, pointing to name1
<code>s = unlink(name)</code>	Remove a directory entry
<code>s = mount(special, name, flag)</code>	Mount a file system
<code>s = umount(special)</code>	Unmount a file system

Agenda

- Recap
- OS Introduction III
 - OS Interface: System Calls
 - OS Internal Structures

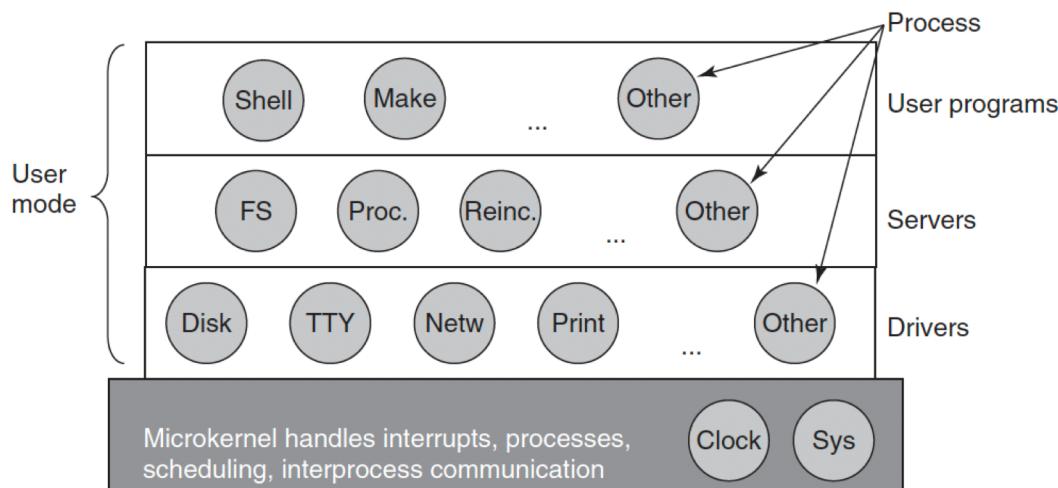
OS Structures

- Monolithic kernel
 - E.g., Linux
 - the entire OS runs as a single program in kernel mode
 - a collection of procedures, linked together into a single large executable
 - may include loadable modules (e.g., drivers, file systems)
 - may include multiple layers logically



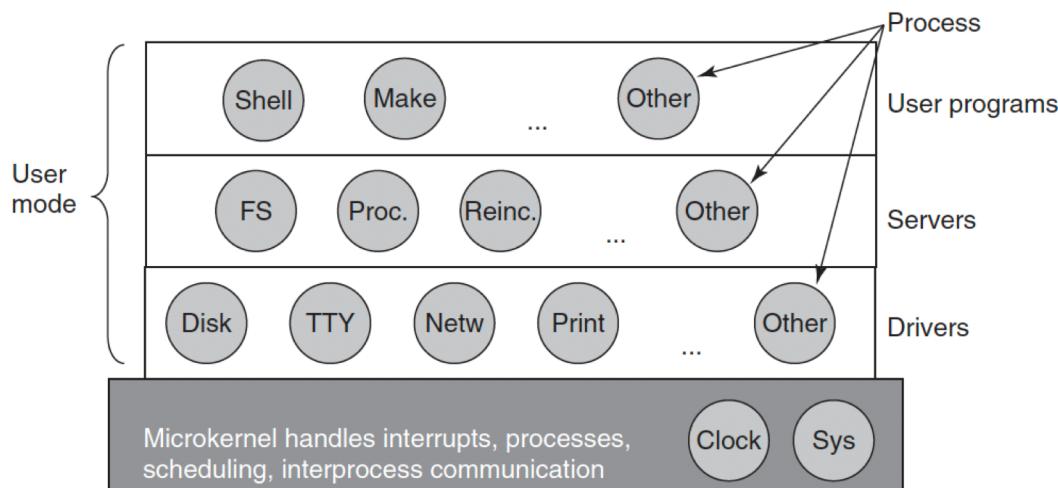
OS Structures

- Microkernel
 - E.g., MINIX 3
 - split the OS up into small, well-defined modules, only one of which—the microkernel—runs in kernel mode
 - the rest run as relatively powerless ordinary user processes



OS Structures

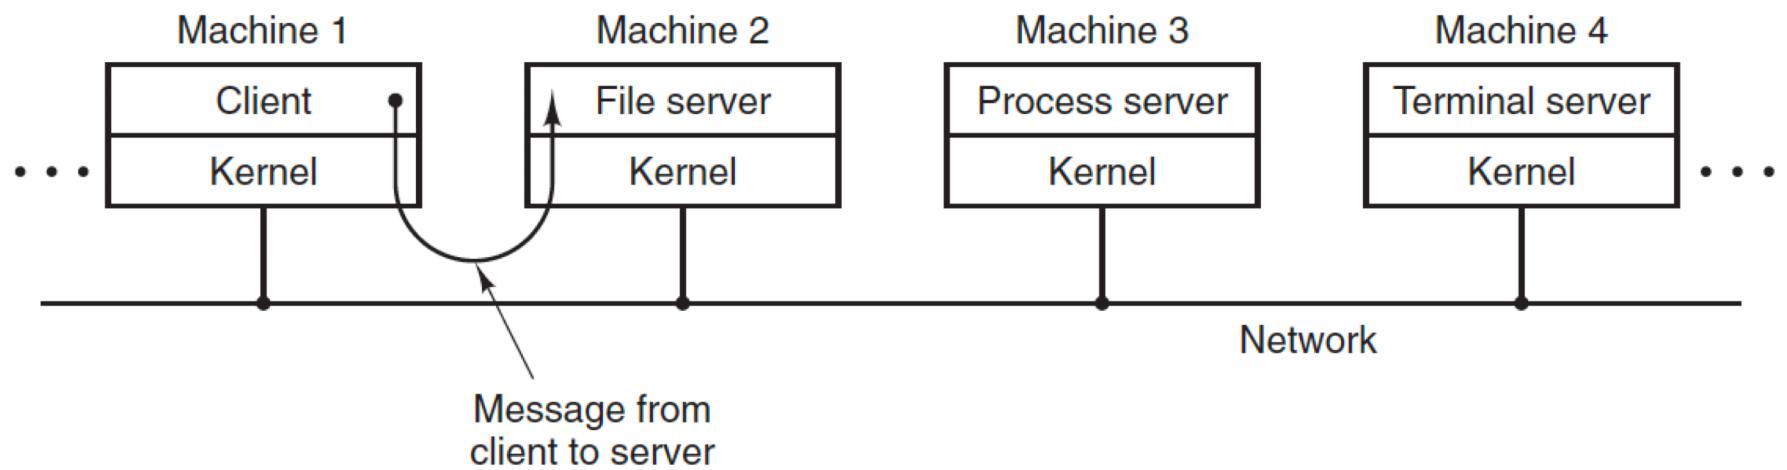
- Microkernel
 - E.g., MINIX 3
 - split the OS up into small, well-defined modules, only one of which—the microkernel—runs in kernel mode
 - the rest run as relatively powerless ordinary user processes



- Advantage/disadvantage compared to Monolithic?

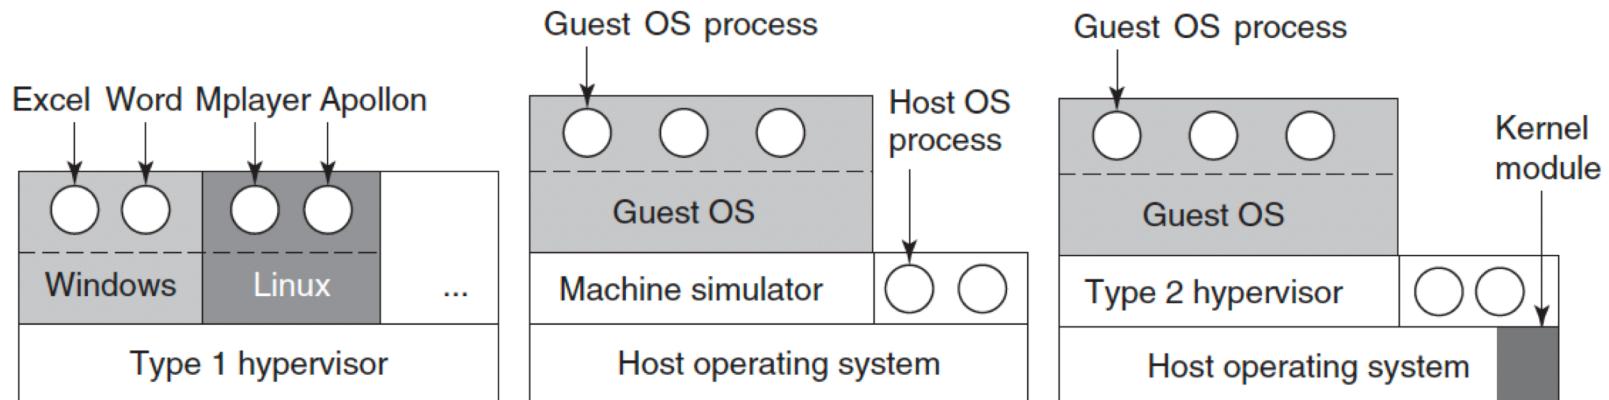
OS Structures

- Client-Server Model
 - distinguish two classes of processes
 - server: each of which provides some service
 - client: uses the services
 - client/server may run on different machines



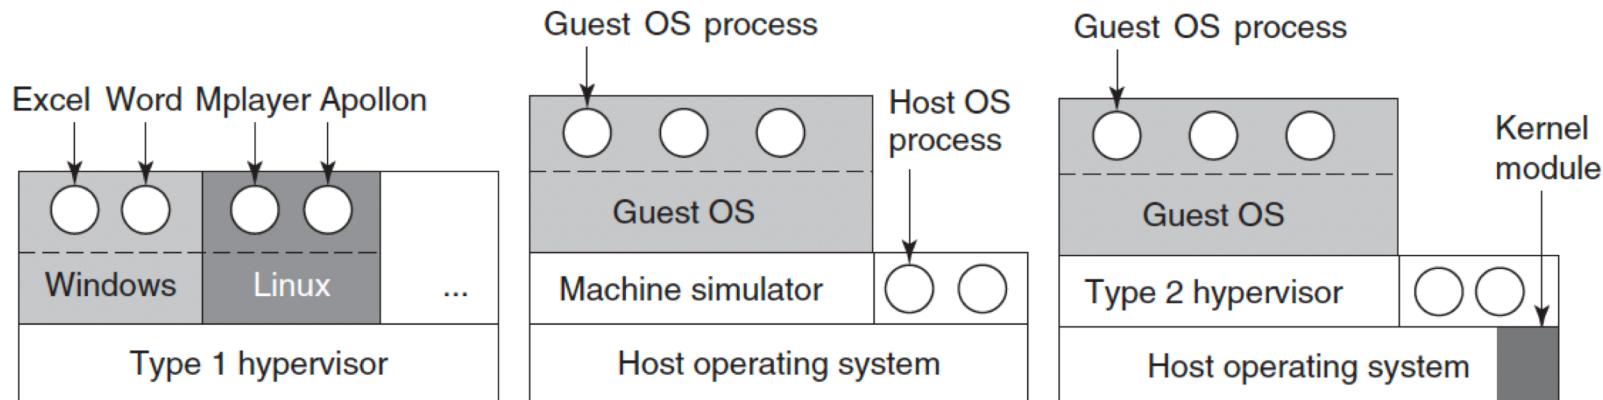
OS Structures

- Virtual Machine (VM)
 - E.g., VMWare, VirtualBox, Xen
 - Another layer of indirection beneath OS
 - Type 1 hypervisor: directly on hardware
 - Type 2 hypervisor: requires support from an underlying OS



OS Structures

- Virtual Machine (VM)
 - E.g., VMWare, VirtualBox, Xen
 - Another layer of indirection beneath OS
 - Type 1 hypervisor: directly on hardware
 - Type 2 hypervisor: requires support from an underlying OS
 - advantage/disadvantage?



Agenda

- Recap
- OS Introduction III
 - OS Interface: System Calls
 - OS Internal Structures

Questions?



*acknowledgement: slides include content from “Modern Operating Systems” by A. Tanenbaum, “Operating Systems Concepts” by A. Silberschatz etc., “Operating Systems: Three Easy Pieces” by R. Arpacı-Dusseau etc., and anonymous pictures from internet.