

IOWA STATE UNIVERSITY

Department of Electrical and Computer Engineering

# Lecture 16: Midterm 1 Review



# Basic Information

- Time
  - 09:00-09:50 am, Oct 4 (Friday)
- Location
  - Marston 2300
- Format
  - Similar to HW1
  - Closed book/notes
- Scope
  - Lecture 1 to today

# Review

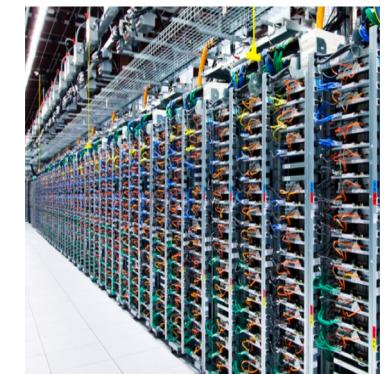
- Introduction
- Process Management
  - Processes
  - Threads
  - Process/Thread Scheduling
  - Inter-Process Communication
  - Classic IPC Problems
  - Deadlocks

# Review

- Introduction
  - Why OS
    - Fundamental to computer systems
      - Affects correctness, security, performance ... of entire system
    - Fundamental to modern society

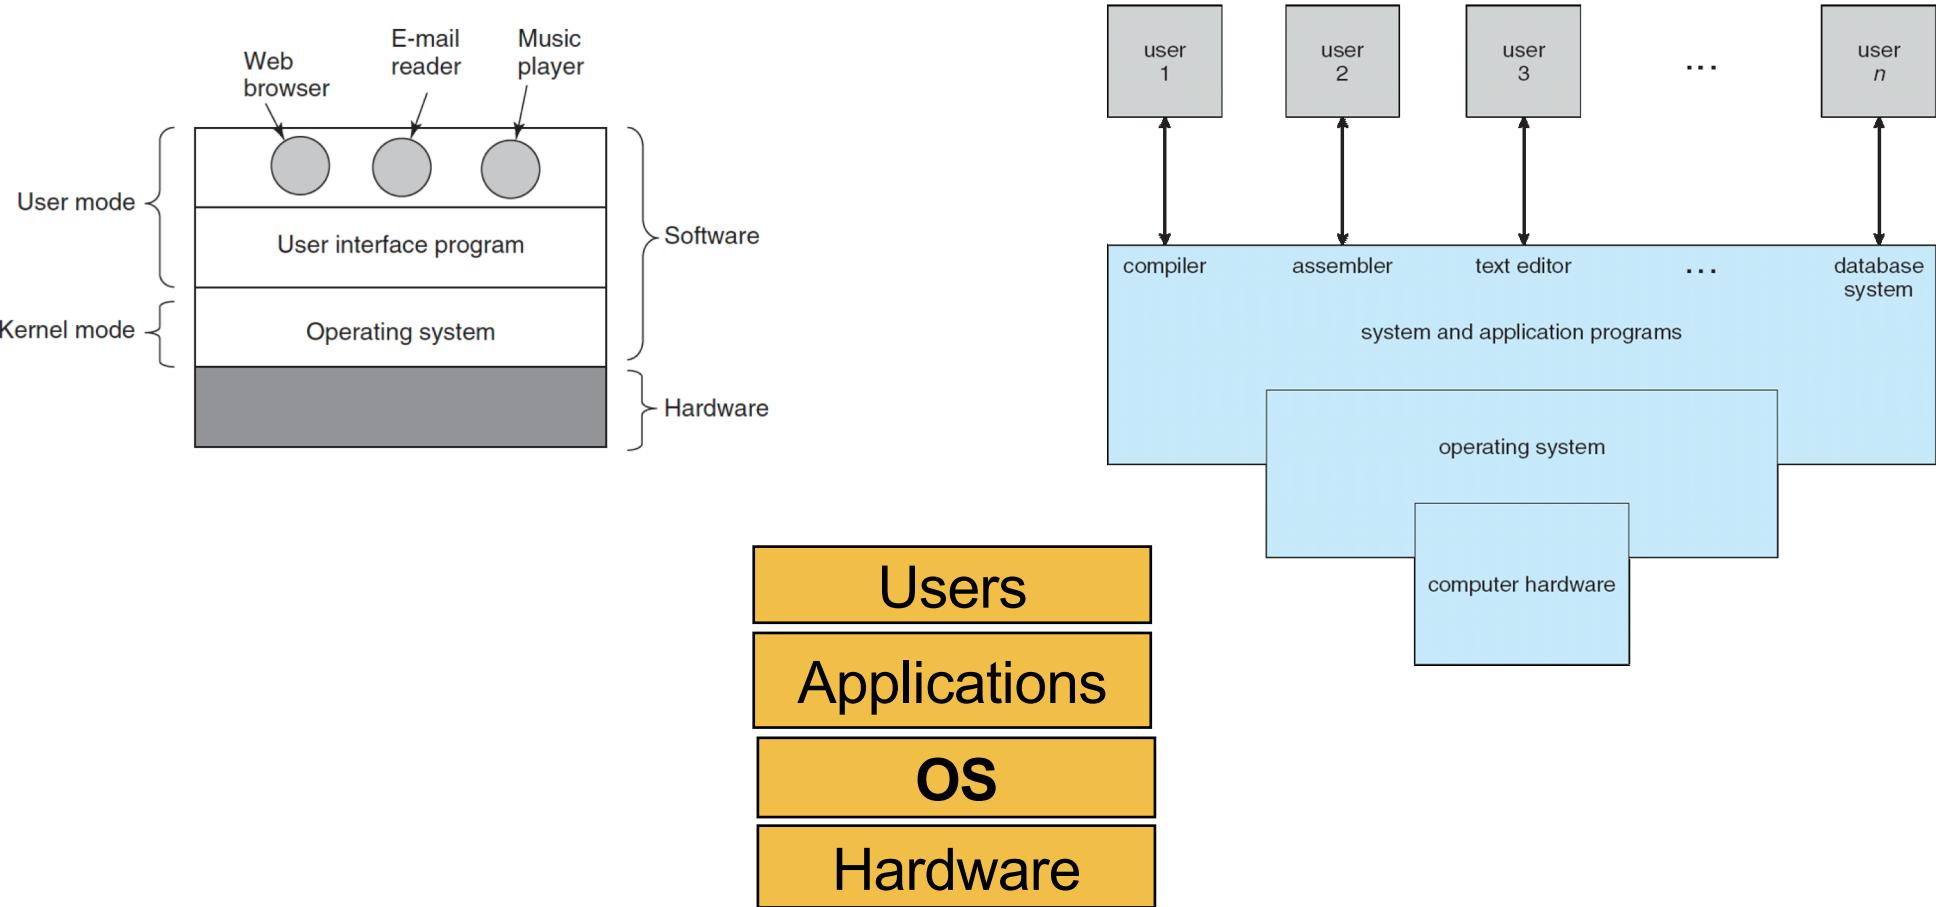


- What's OS
  - OS is a resource manager
  - OS is a control program
    - an extended/virtualized machine with abstraction



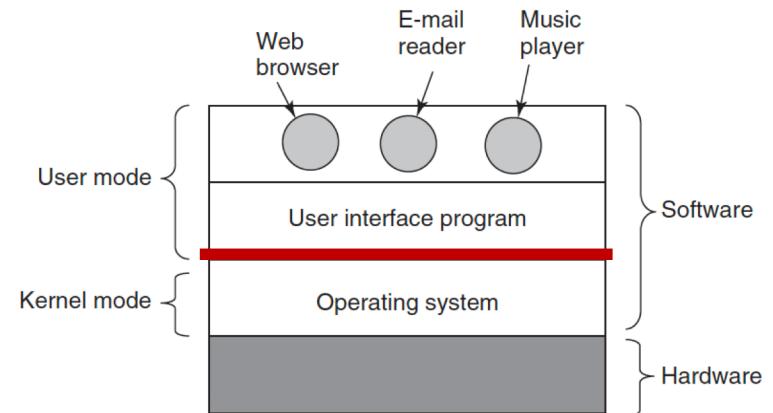
# Review

- Introduction
  - Computer System Structure



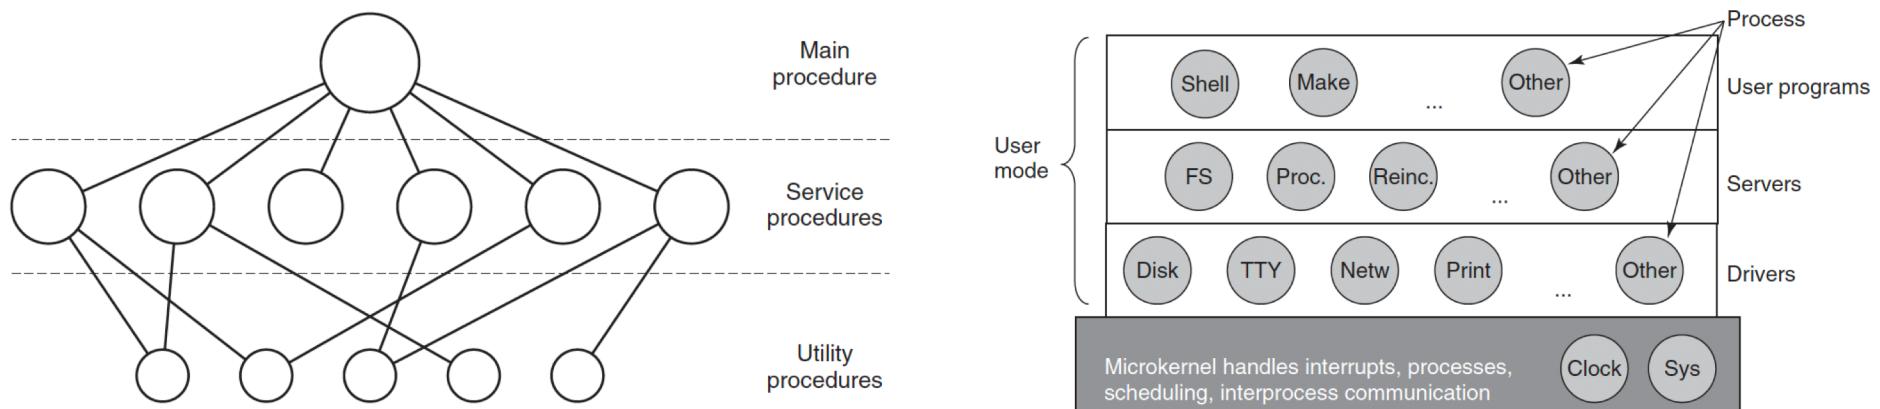
# Review

- Introduction
  - OS Abstractions for HW
    - CPU
      - process and/or thread
    - Memory
      - address space
    - Disks
      - Files
  - System Calls
    - Function calls implemented by OS
    - interface to apps/users



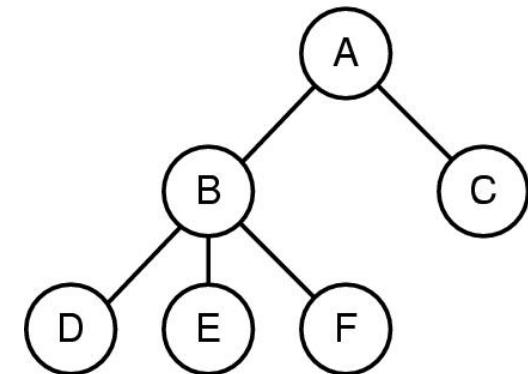
# Review

- Introduction
  - OS Structures
    - Monolithic kernel V.S. Microkernel
      - Monolithic: the OS runs as a single program in kernel mode
      - Microkernel: split the OS up into small, well-defined modules, only one of which—the microkernel—runs in kernel mode



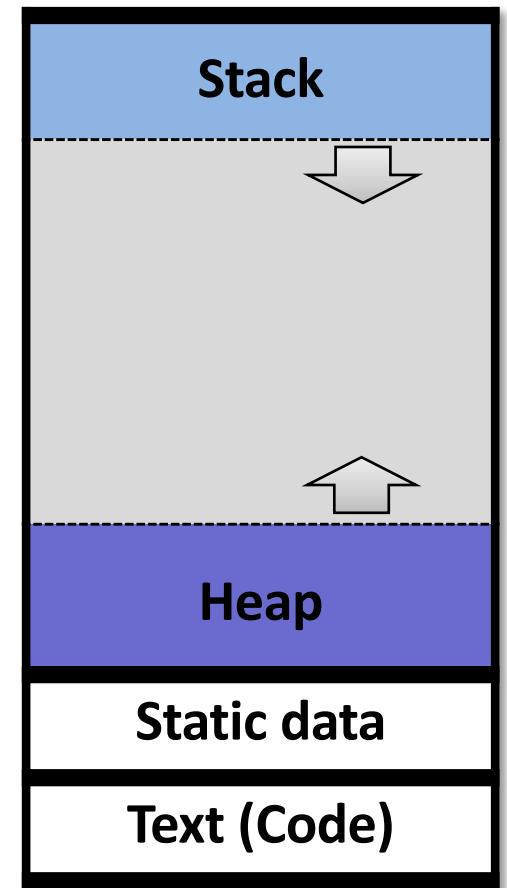
# Review

- Process
  - A program *in execution*
  - holds all the info needed to run a program
    - Address space
      - Program (text), data, stack
    - Register values
      - Program counter, stack pointer, etc
- Process Tree
  - a hierarchy of related processes



# Review

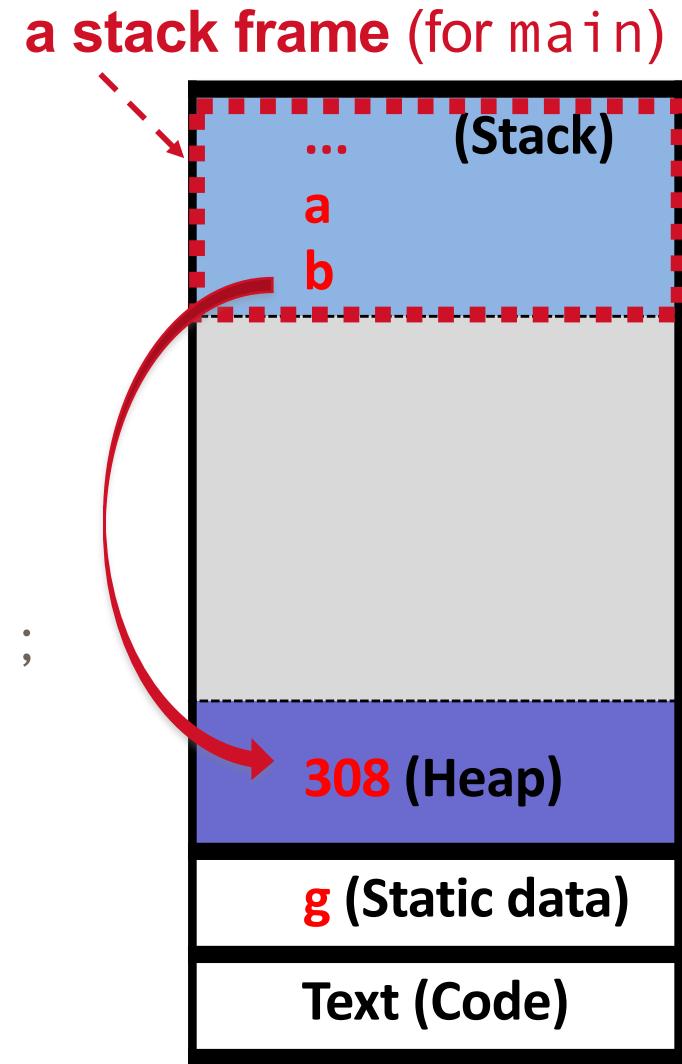
- Process Address Space
  - Four segments
    - stack
      - local variables, function parameters, return address
    - heap
      - dynamically allocated data
        - e.g., malloc()
    - static data
      - global/static variables
    - code (text)
      - instructions of the program



# Review

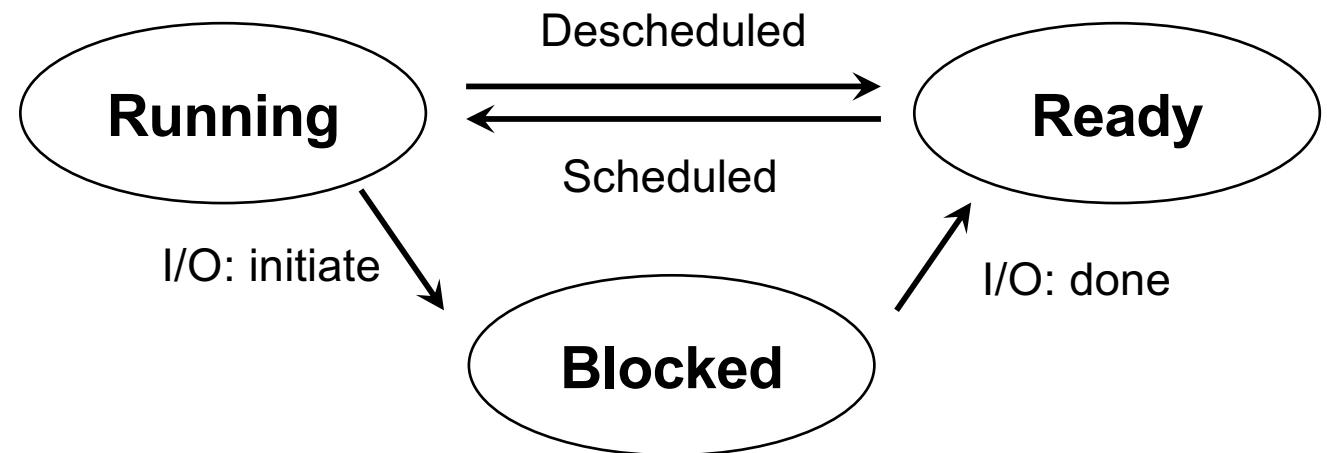
- Process Address Space
  - Example
    - Memory leak

```
int g;  
int main() {  
    int a;  
    int*b = (int*)malloc(sizeof(int));  
    *b = 308;  
    return 0;  
}
```



# Review

- Process States
  - Three basic states
    - Running
    - Ready
    - Blocked
- State transition



# Review

- Process Context
  - Stored in a “Process Control Block (PCB)”

Process management	Memory management	File management
Registers Program counter Program status word Stack pointer Process state Priority Scheduling parameters Process ID Parent process Process group Signals Time when process started CPU time used Children's CPU time Time of next alarm	Pointer to text segment info Pointer to data segment info Pointer to stack segment info	Root directory Working directory File descriptors User ID Group ID

- Context Switch
  - switching the CPU to another process by
    - saving the context of an old process
    - loading the context of a new process

# Review

- Process APIs
  - Example: fork()

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]){
    printf("hello world (pid:%d)\n", (int) getpid()); //get process ID
    int rc = fork();          // create a child process
    if (rc < 0) {            // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {     // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {                 // parent goes down this path (main)
        printf("hello, I am parent of %d (pid:%d)\n",
               rc, (int) getpid());
    }
    return 0;
} c
```

# Review

- Process APIs
  - Example: fork()
    - results (non-deterministic)

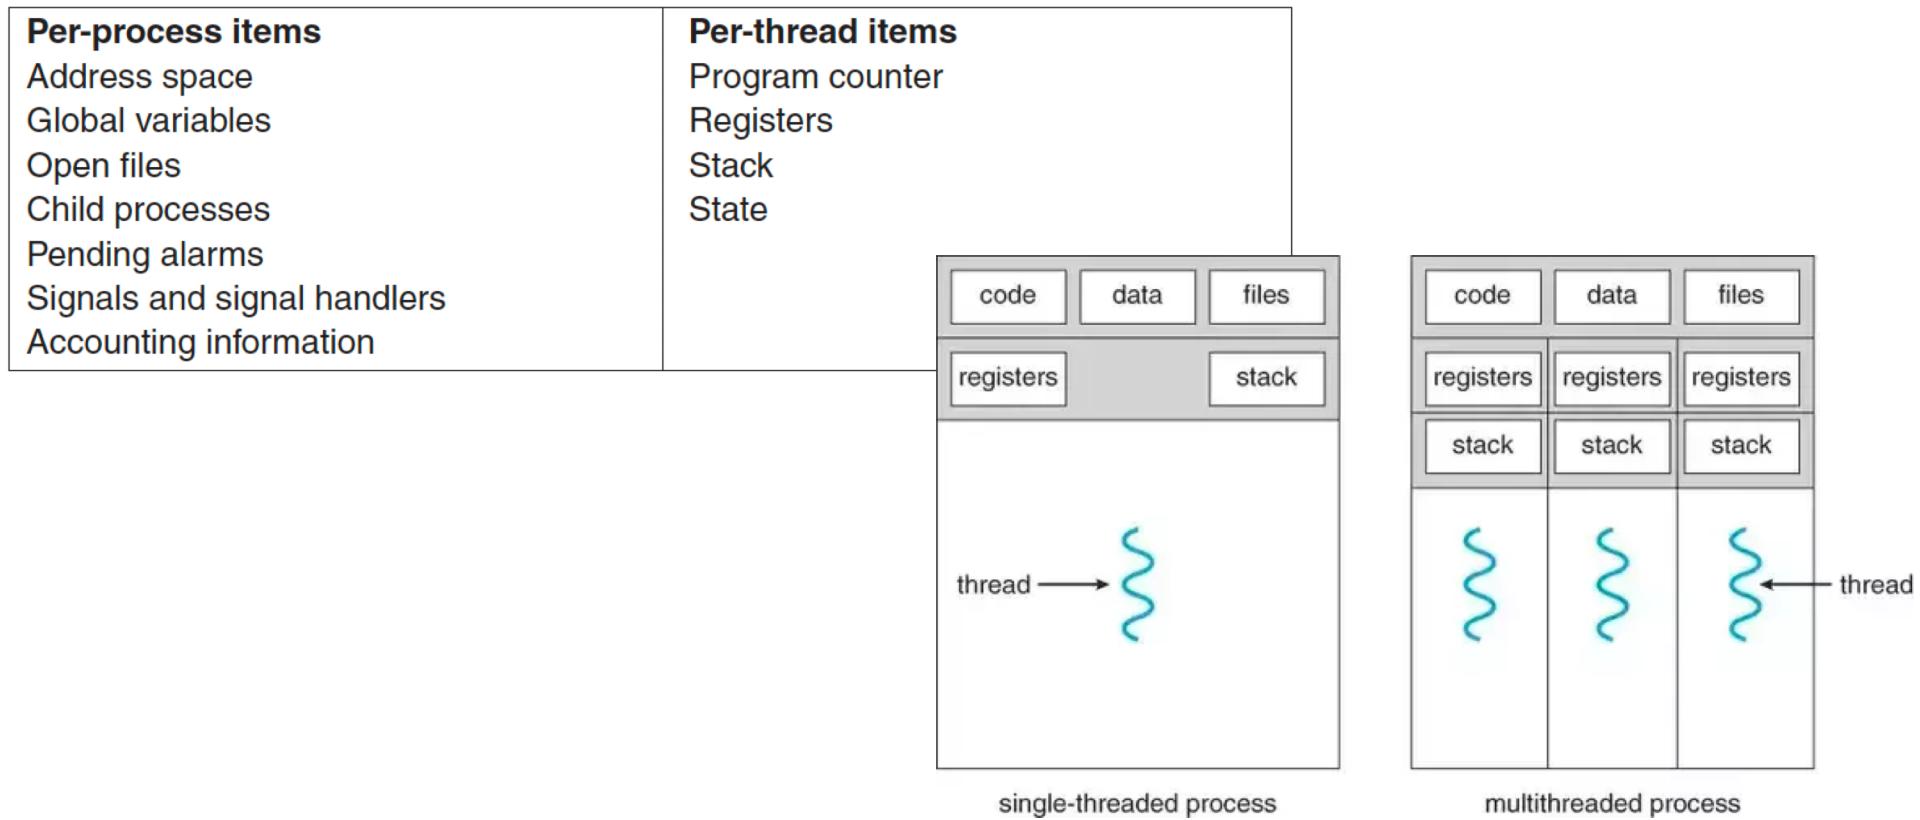
```
prompt> ./p1
hello world (pid:29146)
hello, I am parent of 29147 (pid:29146)
hello, I am child (pid:29147)
prompt>
```

or

```
prompt> ./p1
hello world (pid:29146)
hello, I am child (pid:29147)
hello, I am parent of 29147 (pid:29146)
prompt>
```

# Review

- Multi-threaded Process
  - Multiple threads of control within a process



# Review

- Process/Thread Scheduling
  - First-Come, First-Served (FCFS)
  - Shortest-Job-First (SJF)
  - Shortest Remaining Time Next
  - Round Robin (RR)
  - Priority Scheduling

# Review

- Process/Thread Scheduling
  - First-Come, First-Served (FCFS)
    - Scheduling based on the arrival order of processes

Process	Burst Time
$P_1$	24
$P_2$	3
$P_3$	3

- Assume the processes arrive in the order:  $P_2, P_3, P_1$
- The Gantt Chart for the schedule is:



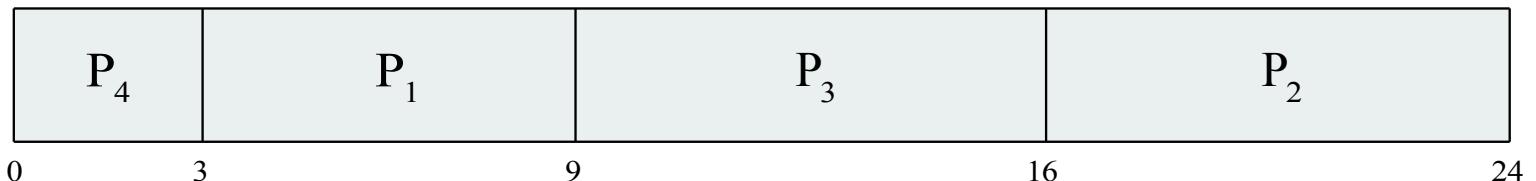
- Waiting time:  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$

# Review

- Process/Thread Scheduling
  - Shortest-Job-First (SJF)
    - Scheduling based on (predicted) CPU burst time

<u>Process</u>	<u>Burst Time</u>
$P_1$	6
$P_2$	8
$P_3$	7
$P_4$	3

- Assume all processes are ready at time 0
- Gantt Chart:



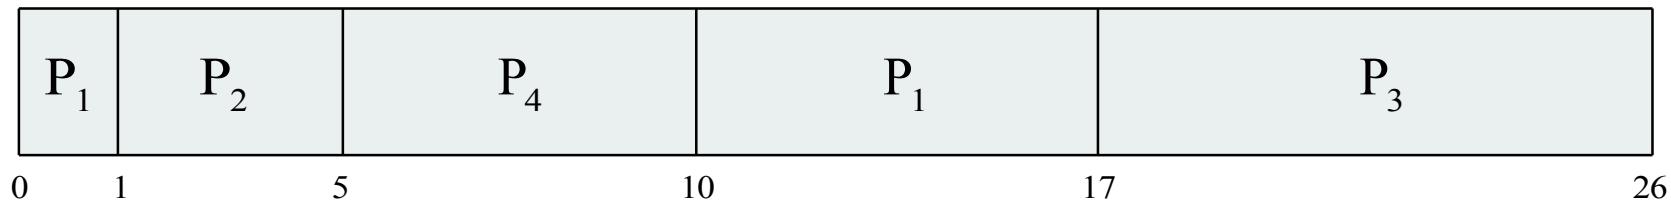
- Average waiting time =  $(3 + 16 + 9 + 0) / 4 = 7$
- Turnaround time: P<sub>1</sub> = 9; P<sub>2</sub> = 24; P<sub>3</sub> = 16; P<sub>4</sub> = 3

# Review

- Process/Thread Scheduling
  - Shortest Remaining Time Next
    - Preemptive; varying arrival time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

- Gantt Chart:



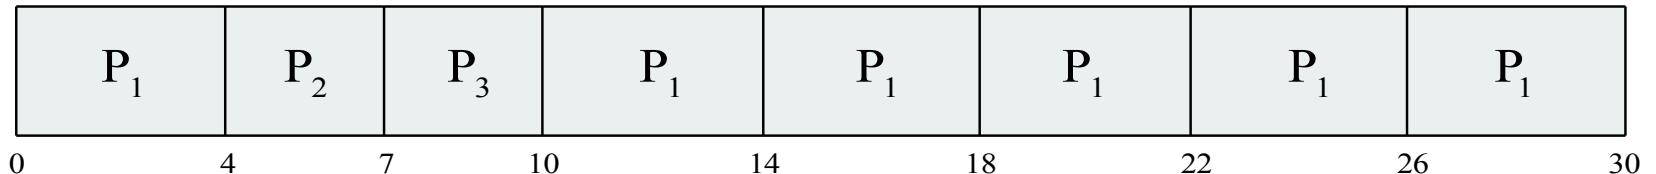
- Average waiting time = 6.5

# Review

- Process/Thread Scheduling
  - Round Robin (RR)
    - Each process is assigned a time interval (quantum) during which it is allowed to run

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Assume quantum  $q = 4$ . The Gantt chart is:



- $q$  should be large compared to context switch time
  - E.g.,  $q$  usually 10ms to 100ms, context switch < 10 usec

# Review

- Process/Thread Scheduling
  - Priority Scheduling
    - schedule the process with the highest priority

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	10	3
$P_2$	1	1
$P_3$	2	4
$P_4$	1	5
$P_5$	5	2

- Assume non-preemptive, all processes are ready at time 0
- Gantt Chart:



- Average waiting time =  $(1+6+16+18)/5 = 8.2$

# Review

- Inter-Process Communications (IPC)
  - Two basic methods
    - Shared memory
    - Message passing
  - Race Condition (Data Race)
    - two or more processes are reading or writing some **shared data** and the final result depends on who runs precisely when
  - Critical Region
    - A piece of code that **accesses a shared variable** and must not be concurrently executed by more than one thread
      - Multiple threads executing critical section can result in a race condition.
      - Need to support **mutual exclusion**

# Review

- Example of race condition
  - Two threads perform “counter = counter +1”
    - “counter” is a shared variable; initially counter = 50

OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
			100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt					
	save T1's state				
	restore T2's state		100	0	50
		mov 0x8049a1c, %eax	105	50	50
		add \$0x1, %eax	108	51	50
		mov %eax, 0x8049a1c	113	51	51
interrupt					
	save T2's state				
	restore T1's state		108	51	50
		mov %eax, 0x8049a1c	113	51	51

# Review

- Solutions of mutual exclusion
  - Software solution
    - Disabling interrupts
      - Single processor only
      - Use for kernel
    - Strict alternation
      - Strict ordering
      - Busy waiting
    - Peterson's solution
      - Busy waiting
  - Hardware solution
    - TSL/XCHG
      - Work on multiprocessors
      - Busy waiting



# Review

- Solutions of mutual exclusion
  - Pthread locks

```
pthread_mutex_t lock;  
...  
pthread_mutex_lock(&lock);  
x = x + 1; // or whatever your critical section is  
pthread_mutex_unlock(&lock);
```



# Review

- Semaphore
  - Semaphore **S** – integer variable
    - Can only be accessed via two indivisible (atomic) operations
      - down() and up()
      - originally called **P()** and **V()**
      - also called **wait()** and **signal()**
    - meaning of **down()** operation (atomic):

```
down(S) {  
    while (S <= 0)  
        ; // busy waiting  
    S--;  
}
```
    - meaning of **up()** operation (atomic):

```
up(S) {  
    S++;  
}
```

# Review

- Semaphore Example
  - Readers-Writers Problem

READER:

```
While (1) {
    down(protector);
    rc++;
    if (rc == 1) //first reader
        down(database);
    up(protector);

    read();

    down(protector);
    rc--;
    If (rc == 0) then // last one
        up(database);
    up(protector);
    ....
}
```

WRITER:

```
While (1) {
    generate_data();
    down(database);
    write();
    up(database);
}
```

Two semaphores:

database  
protector

Initial: protector=1, database =1  
rc =0

# Review

- Conditional Variable
  - Allows a thread to wait till a condition is satisfied
  - Example:

```
int    thread1_done = 0;  
  
pthread_cond_t  cv;  
pthread_mutex_t mutex;
```

Thread 1:

```
printf("hello ");  
pthread_mutex_lock(&mutex);  
thread1_done = 1;  
pthread_cond_signal(&cv);  
pthread_mutex_unlock(&mutex);
```

Thread 2:

```
pthread_mutex_lock(&mutex);  
while (thread1_done == 0) {  
    pthread_cond_wait(&cv, &mutex);  
}  
printf(" world\n");  
pthread_mutex_unlock(&mutex);
```

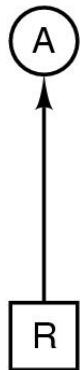
# Review

- Deadlocks
  - Four conditions need to hold for a deadlock to occur

Condition	Description
Mutual Exclusion	Threads claim exclusive control of resources that they require.
Hold-and-wait	Threads hold resources allocated to them while waiting for additional resources
No preemption	Resources cannot be forcibly removed from threads that are holding them.
Circular wait	There exists a circular chain of threads such that each thread holds one more resources that are being requested by the next thread in the chain

# Review

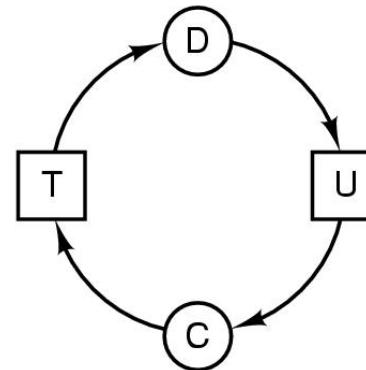
- Deadlocks
  - Resource allocation graph
    - resource R assigned to process A
    - process B is requesting/waiting for resource S
    - process C and D are in deadlock over resources T and U



(a)



(b)



(c)

# Agenda

- Midterm 1 Review

## Questions?



\*acknowledgement: slides include content from “Modern Operating Systems” by A. Tanenbaum, “Operating Systems Concepts” by A. Silberschatz etc., “Operating Systems: Three Easy Pieces” by R. Arpaci-Dusseau etc., and anonymous pictures from internet.