

IOWA STATE UNIVERSITY

Department of Electrical and Computer Engineering

Lecture 18: Segmentation & Free-Space Management

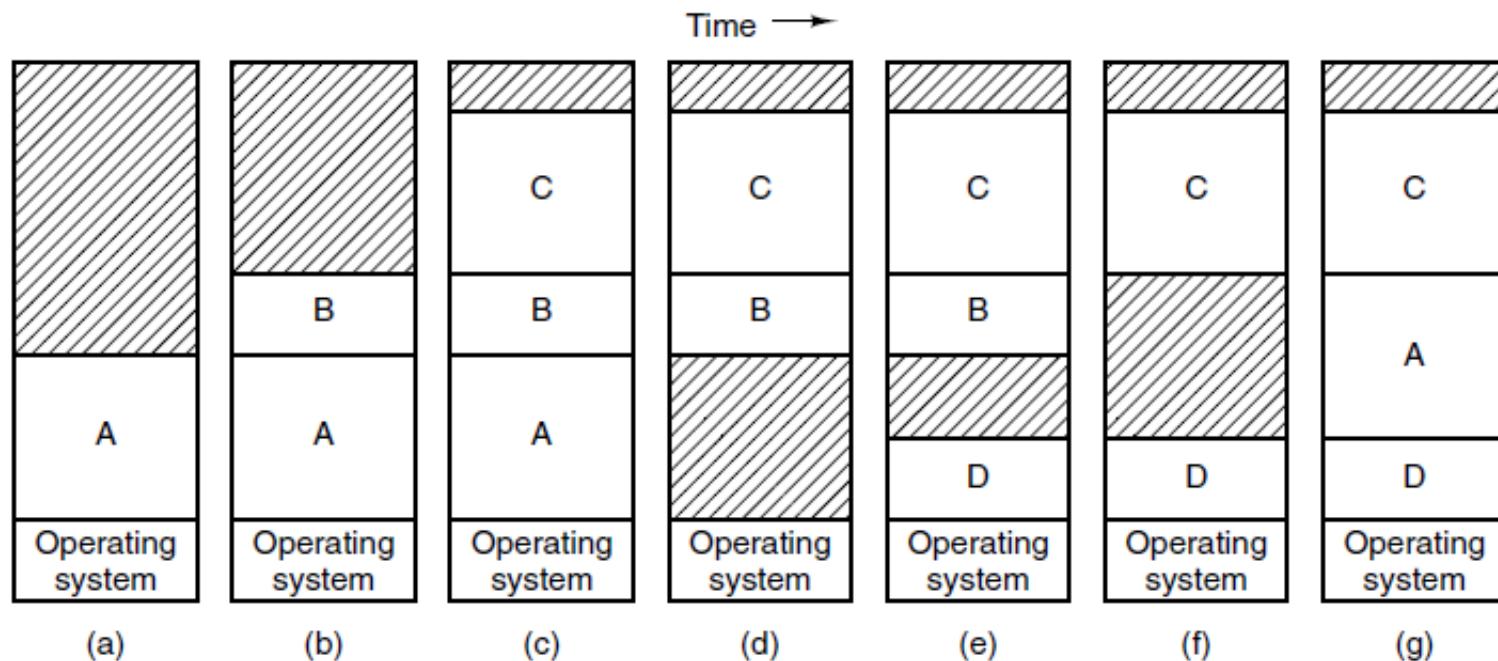


Agenda

- Recap
- Segmentation
- Free-Space Management

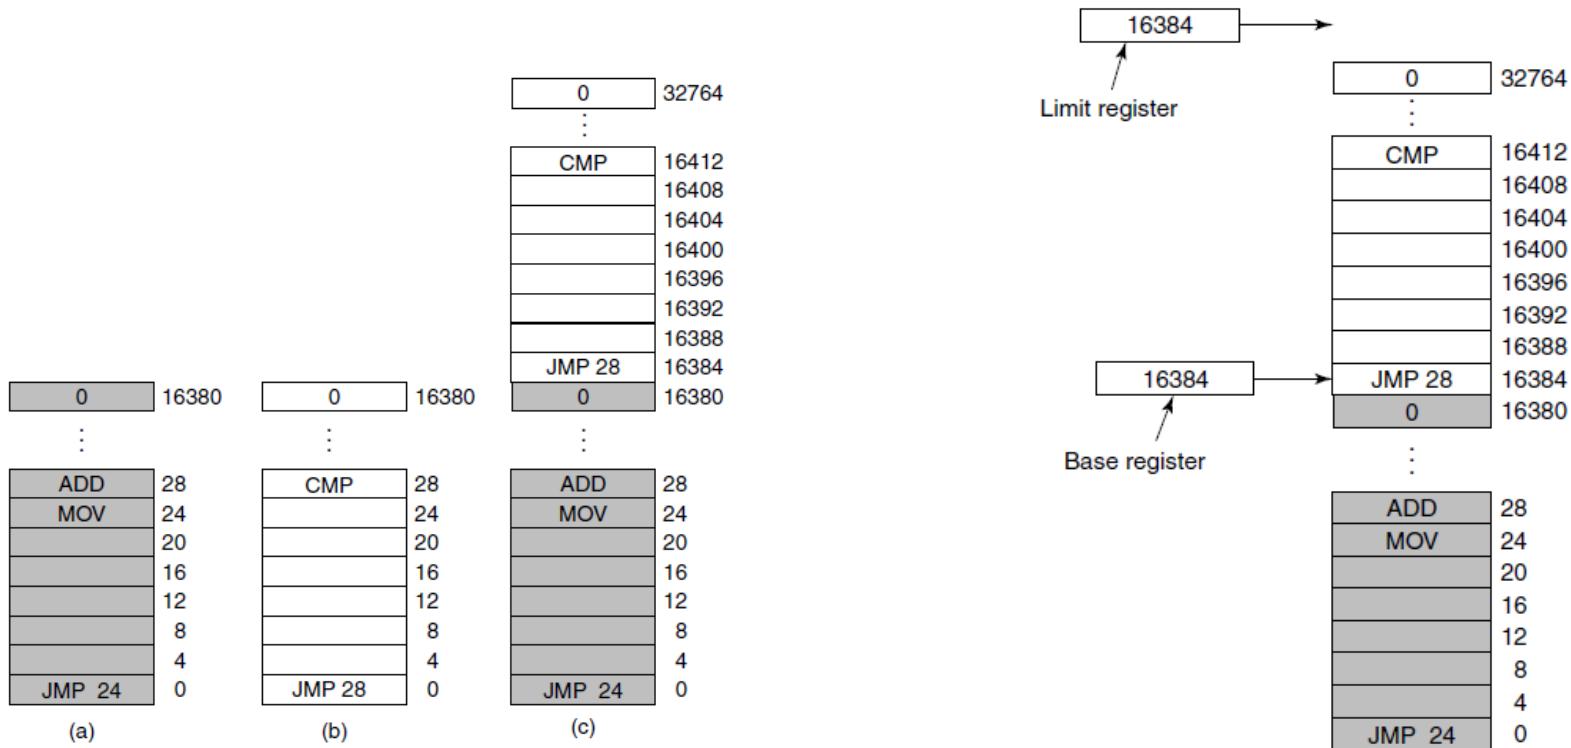
Recap

- Memory management in early systems
 - Treat the memory space of a process as a whole



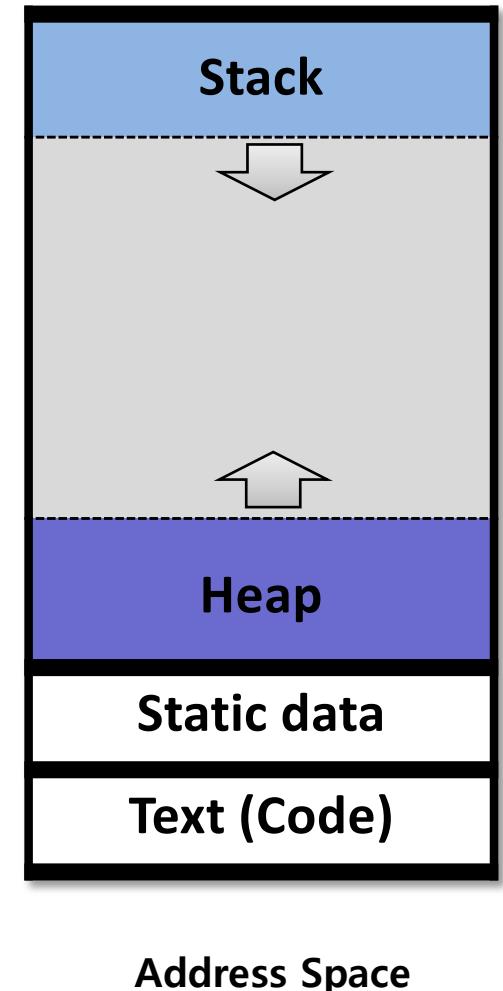
Recap

- Memory management in early systems
 - Relocation using a pair of *base* and *limit* registers
 - Basic protection: checking if address within the range



Recap

- Modern OS provides a virtual memory for each process
 - An illusion that each process uses the whole memory itself
 - Improve memory efficiency, isolation & protection
- Address space
 - An abstraction of physical memory for a process
 - the set of all virtual addresses visible to a program

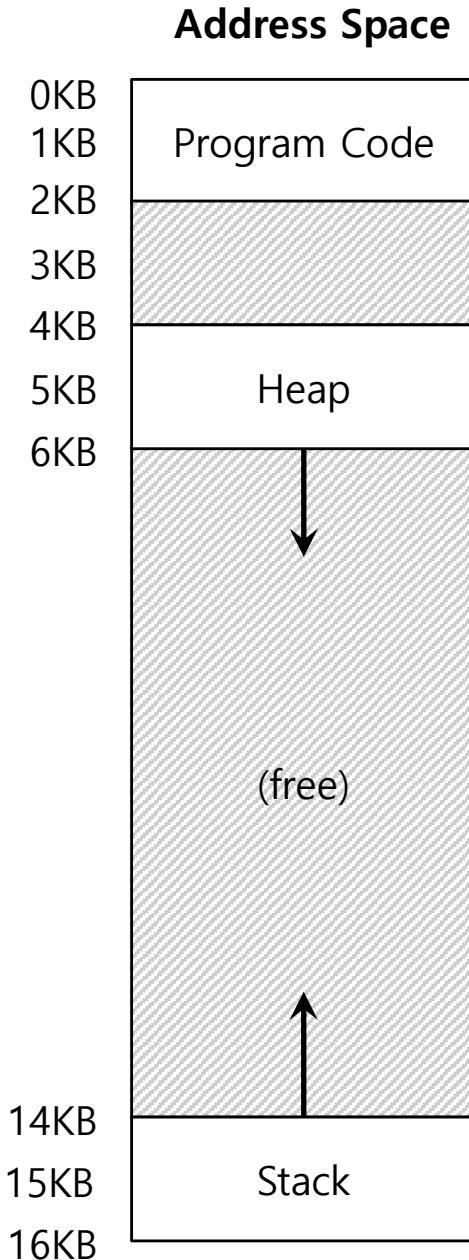


Agenda

- Recap
- Segmentation
- Free-Space Management

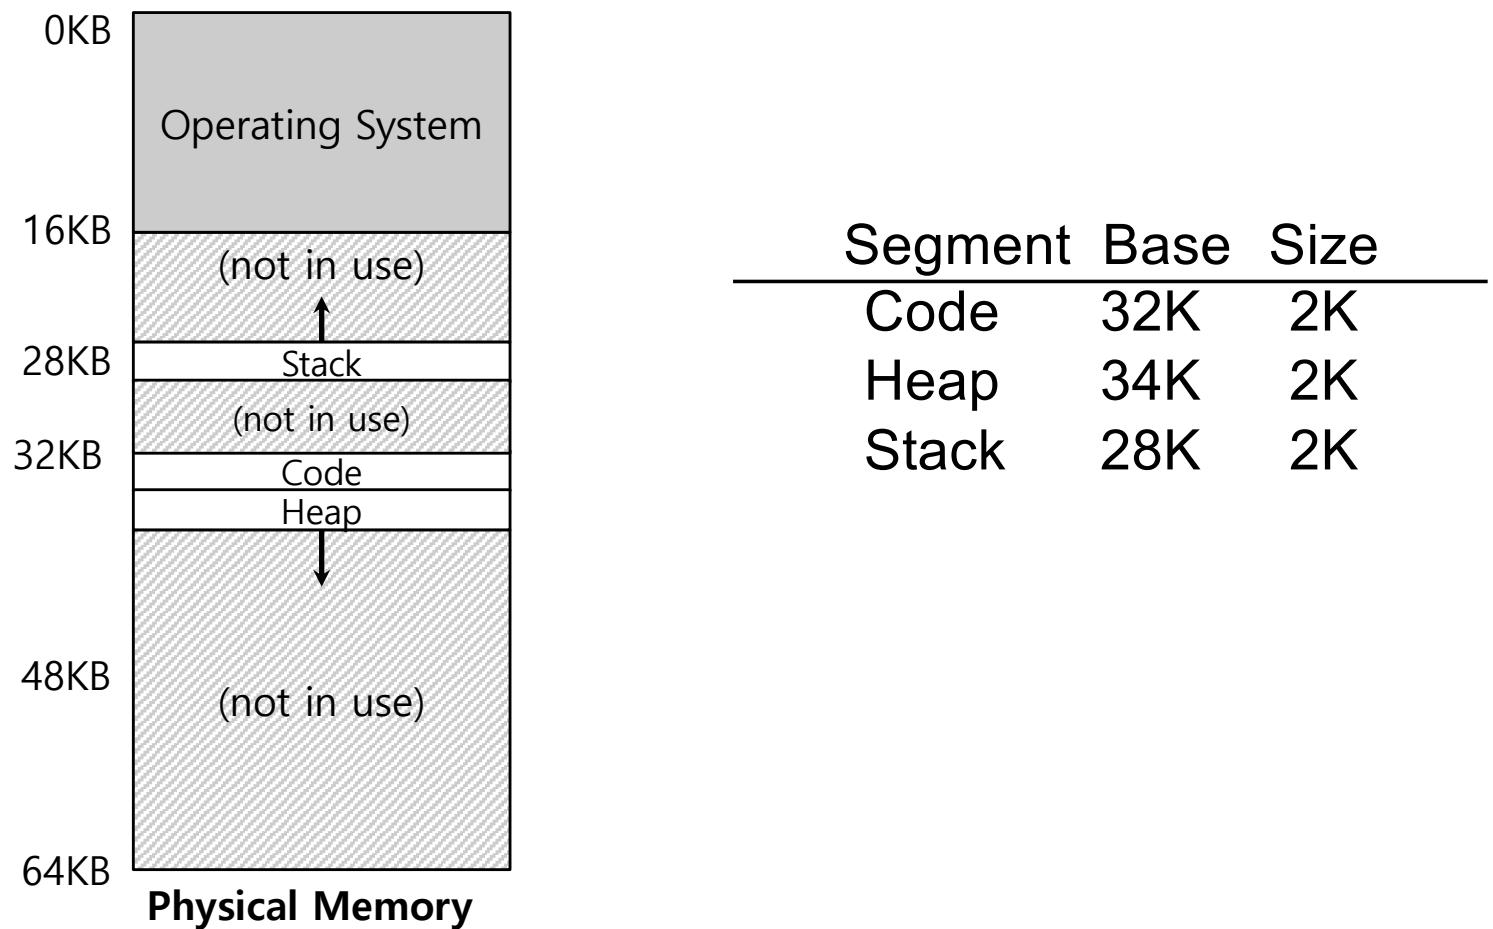
Segmentation

- Break the full address space into a few segments
 - Segment: a contiguous portion of the address space of a particular length
 - Logically-different segment:
 - code, stack, heap, data
 - Each segment can be placed in different part of physical memory
 - Base and limit exist for each segment



Segmentation

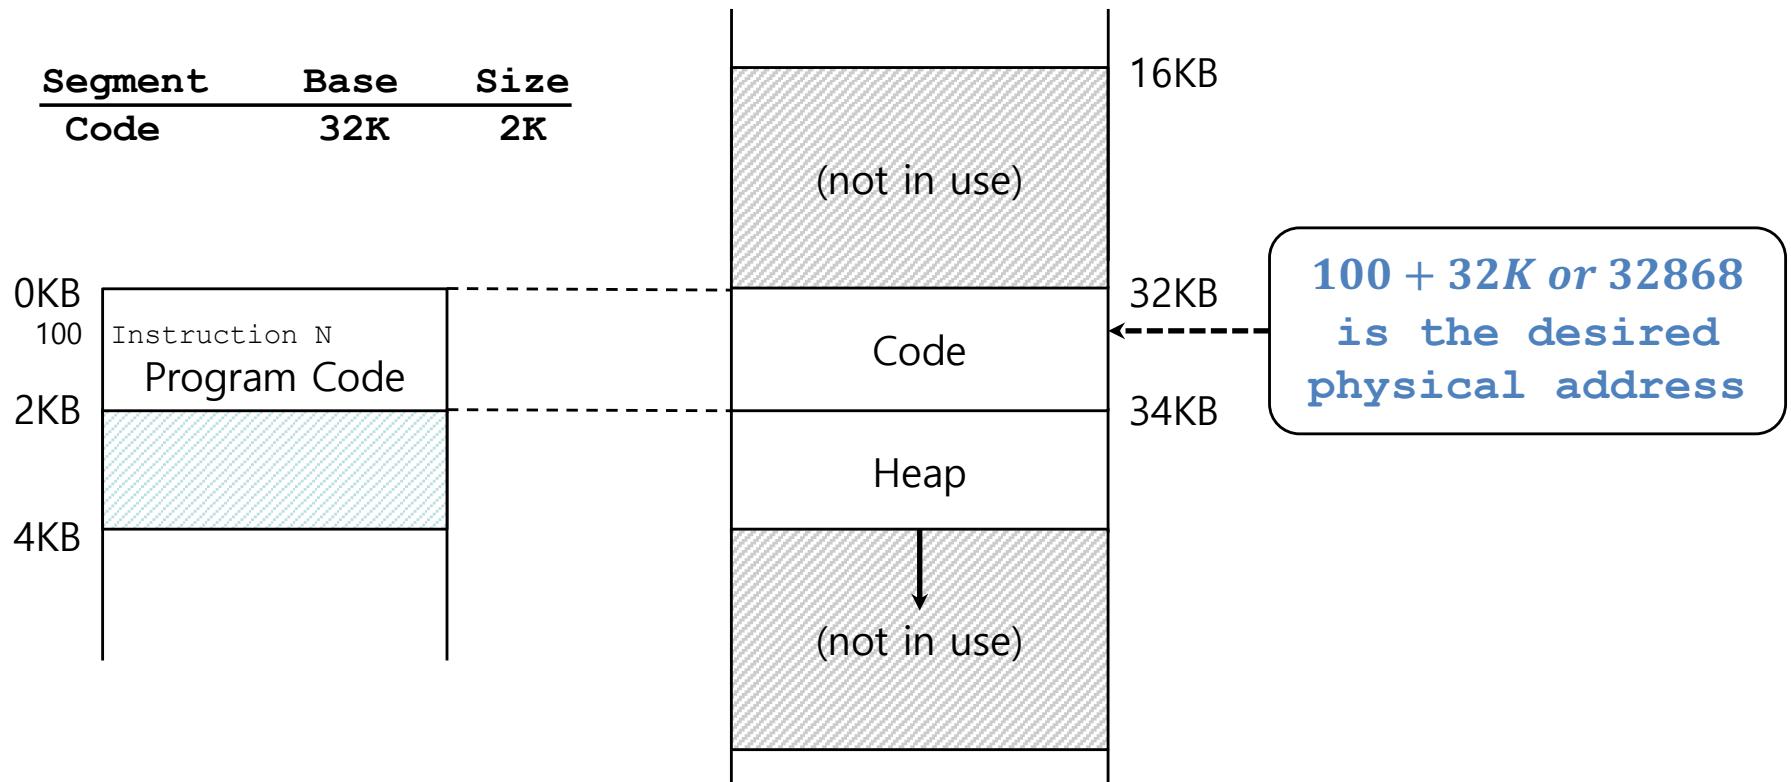
- Placing segments in physical memory



Address Translation of Segmentation

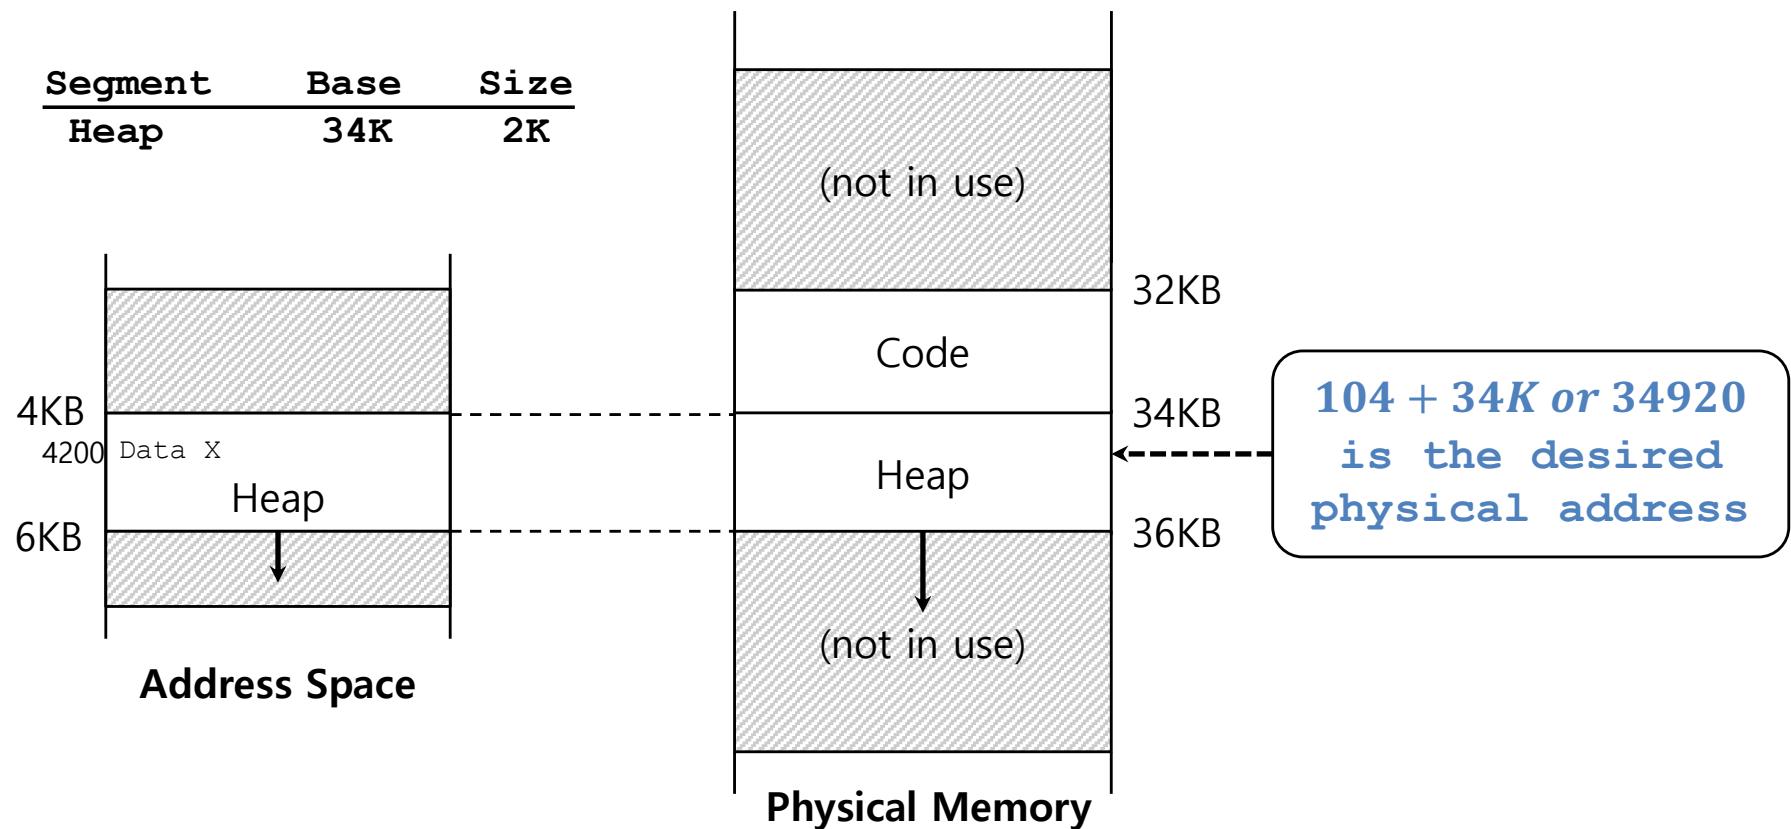
$$\text{physical address} = \text{offset} + \text{base}$$

- The code segment starts at virtual address 0 in address space
 - The offset of virtual address 100 in code region is 100.



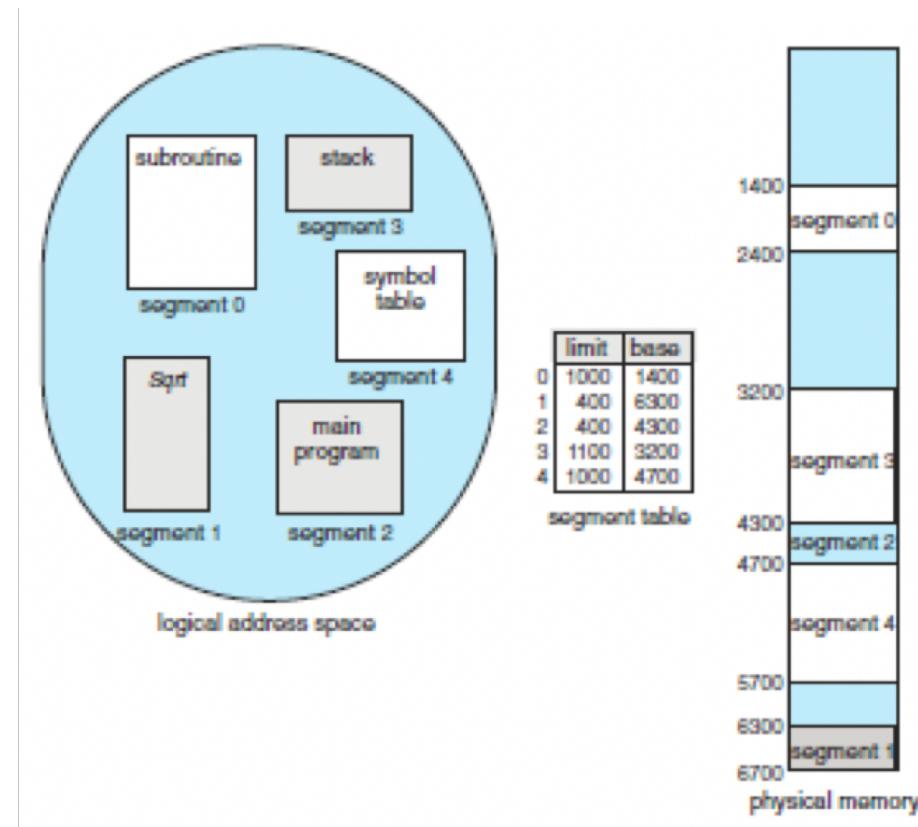
Address Translation of Segmentation

- *Virtual address+base* is not always the correct physical address.
 - The heap segment starts at virtual address 4K in address space.
 - The offset of virtual address 4200 is 104.



Address Translation of Segmentation

- To support multiple segments, a segment table is needed

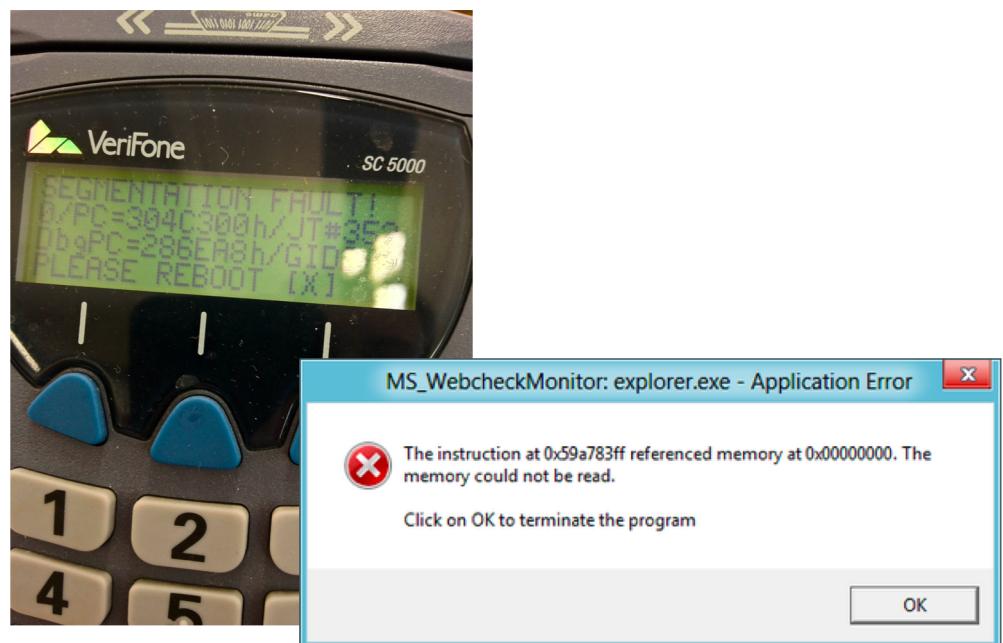


Segmentation Fault

- In case of illegal address
 - Hardware detects the illegal access and raises an exception
 - CPU executes a fault handler (part of the OS)
 - E.g., kill the process and throws a *segmentation fault* message to user

```
diego@cryptos:/tmp$ cat segfault.c
int main(int argc, const char *argv[]){
    int *ptr;
    *(ptr + 1000000) = 10;

    return 0;
}
diego@cryptos:/tmp$ gcc segfault.c -o segfault
diego@cryptos:/tmp$ ./segfault
Segmentation fault (core dumped)
diego@cryptos:/tmp$
```



Sharing & Protection

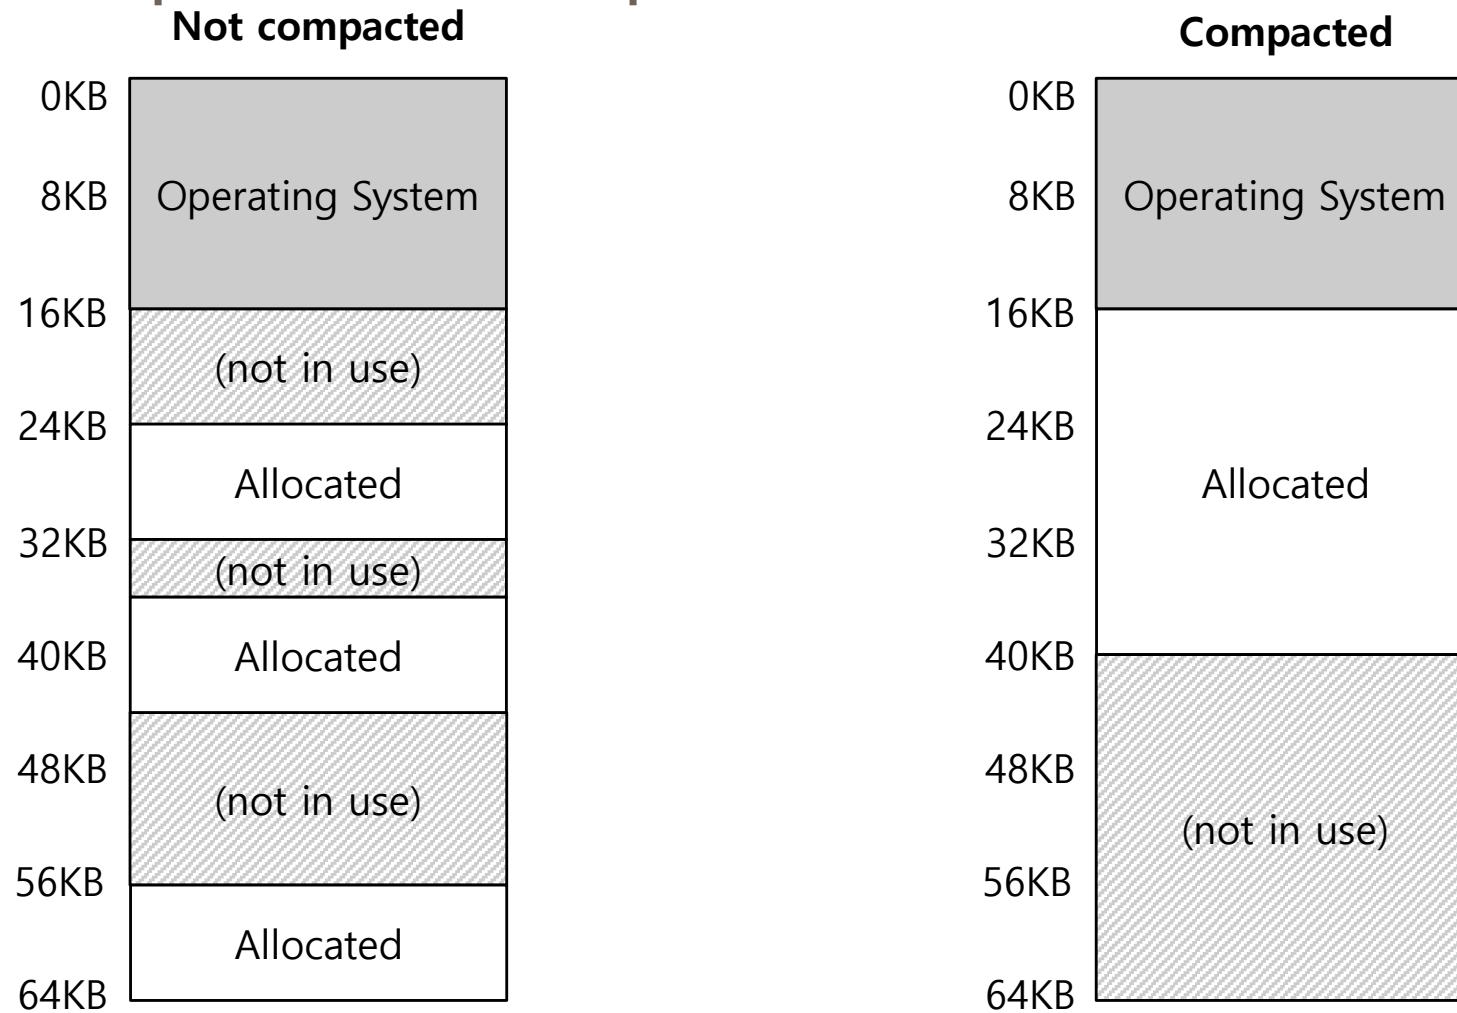
- Segment can be shared between address spaces
 - Code sharing is still in use in systems today
- Extra hardware support is needed for protection
 - A few more bits in the segment registers to indicate permissions of read, write and execute

Fragmentation & Compaction

- External Fragmentation
 - little holes of free space in physical memory that make difficulty to allocate new segments
 - E.g., there is 24KB free space, but not in one contiguous segment.
 - the OS cannot satisfy a 20KB request.
- Compaction
 - rearranging the exiting segments in physical memory.
 - Compaction is costly.
 - Stop running process.
 - Copy data to somewhere.
 - Change segment register value.

Fragmentation & Compaction

- Compaction example



Agenda

- Recap
- Segmentation
- Free-Space Management

Free-Space Management

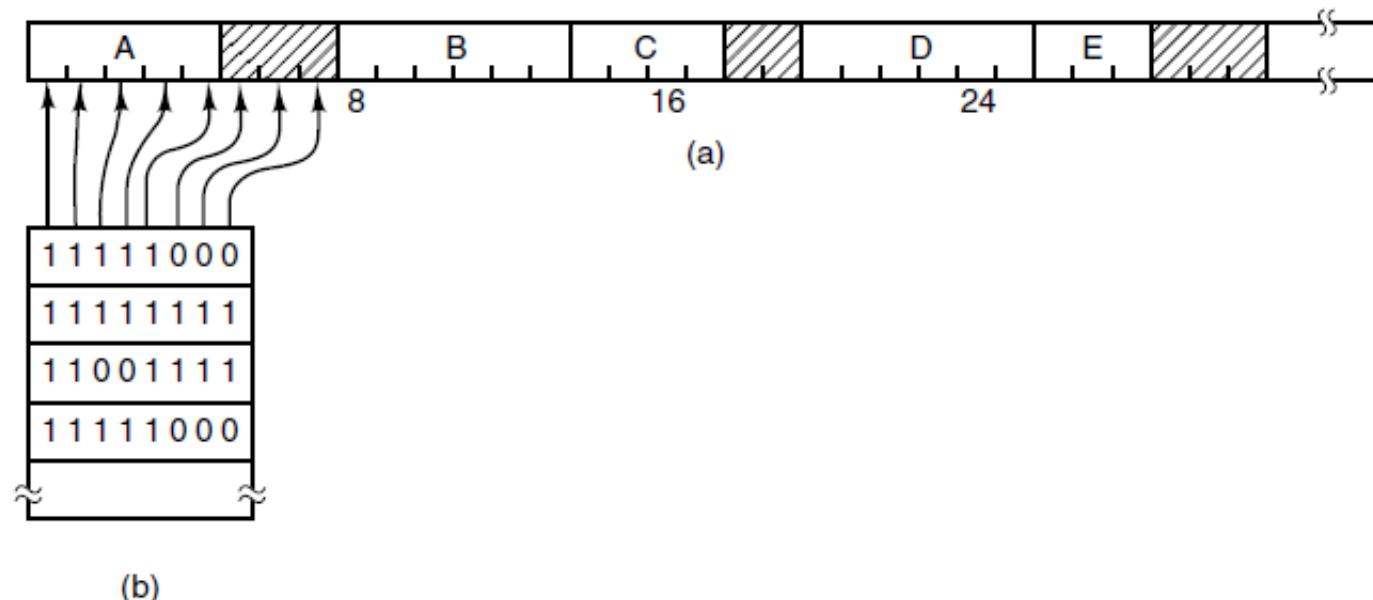
- How to find a free chunk of memory that can satisfy the user request?

Free-Space Management

- How to find a free chunk of memory that can satisfy the user request?
 - Need to keep track of the free/used space
 - Often needs to split a free chunk into two
 - one sub-chunk is allocated to the requesting process
 - the rest of the chunk remains free
- Two Basic Approaches
 - Bitmap
 - Linked list

Free-Space Management

- Bitmap
 - Divide memory into blocks
 - Use a bit 0/1 for every block to indicate whether it is free or not



Free-Space Management

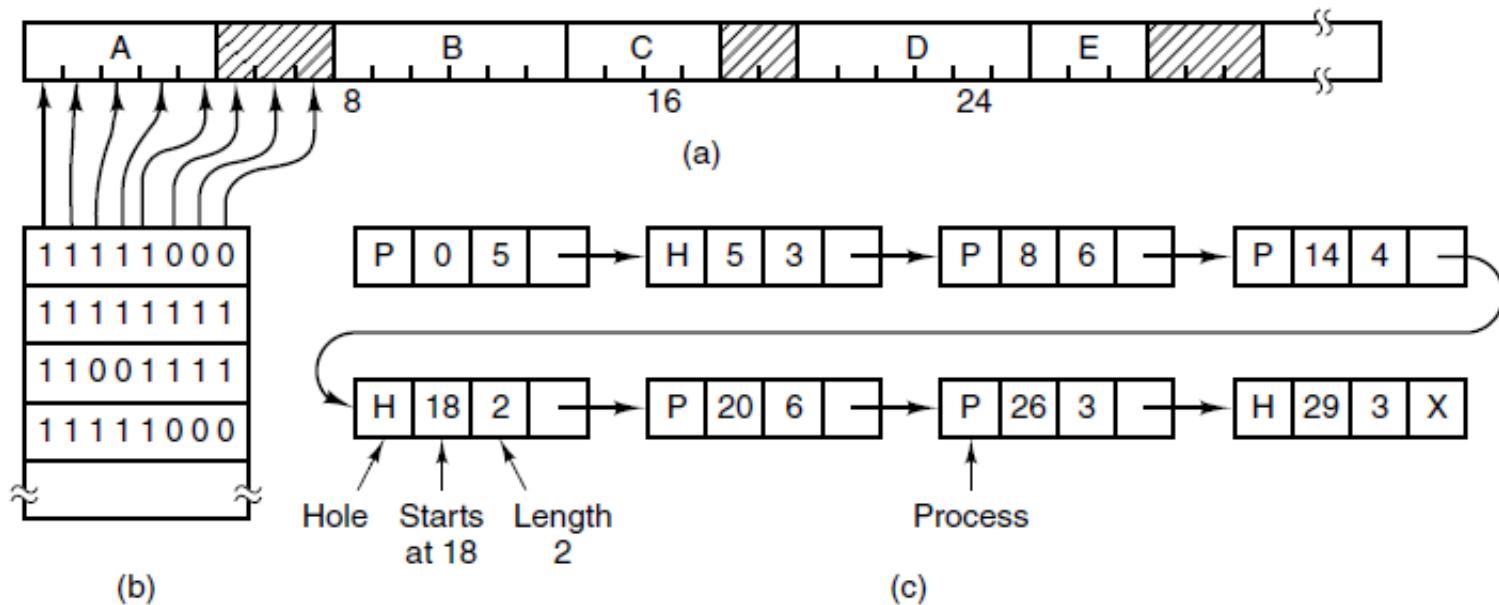
- Bitmap
 - Memory allocation
 - Suppose a process needs k blocks
 - Find k consecutive 0 bits in the bitmap
 - Memory deallocation
 - Reset the corresponding blocks in bitmap to 0

Free-Space Management

- Linked list
 - Version 1:
 - Put all memory blocks in a linked list
 - Each node in this list has four fields
 - First field: indicate free or not
 - Second field: indicate the start address
 - Third field: length of the memory segment
 - Fourth field: a pointer points to the next node

Free-Space Management

- Linked list
 - Version 1: e.g.



Agenda

- Recap
- Segmentation
- Free-Space Management

Questions?



*acknowledgement: slides include content from “Modern Operating Systems” by A. Tanenbaum, “Operating Systems Concepts” by A. Silberschatz etc., “Operating Systems: Three Easy Pieces” by R. Arpacı-Dusseau etc., and anonymous pictures from internet.